

End to End Machine Learning Workflows

Animesh Singh
@AnimeshSingh
STSM and Chief Architect
Data and AI Open Source Platform

IBM Developer



Center for Open Source Data and AI Technologies (CODAIT)

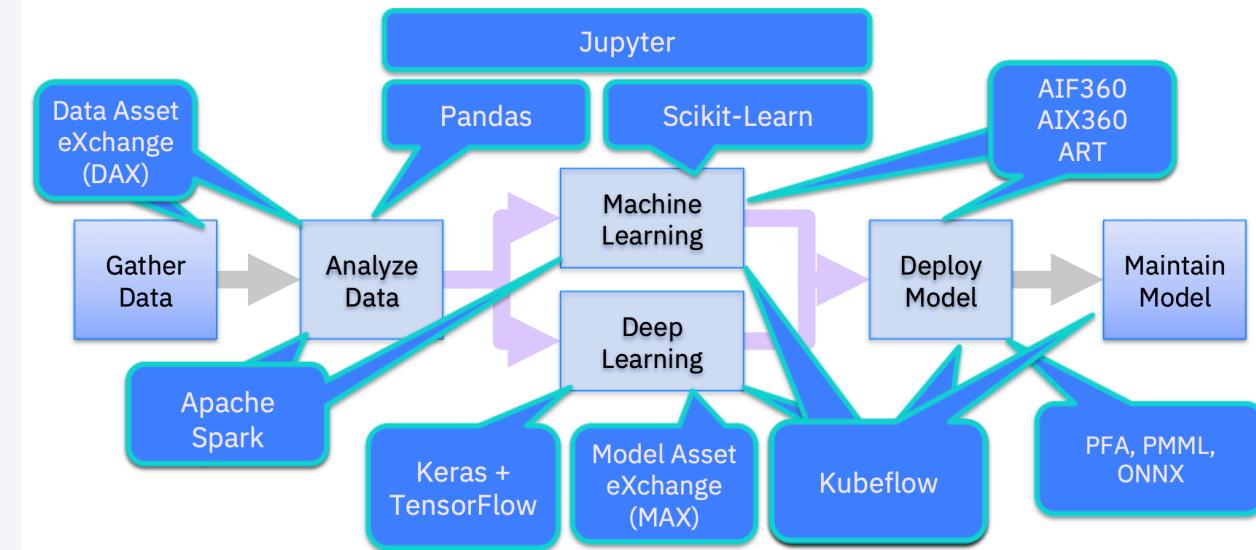
Code – Build and improve practical frameworks to enable more developers to realize immediate value.

Content – Showcase solutions for complex and real-world AI problems.

Community – Bring developers and data scientists to engage with IBM

- Team contributes to over **10 open source projects**
- **17 committers** and many contributors in Apache projects
- Over **1100 JIRAs** and **66,000 lines of code** committed to Apache Spark itself; over **65,000 LoC** into SystemML
- Over **25 product lines** within IBM leveraging Apache Spark
- Speakers at over 100 conferences, meetups, unconferences and more

Improving Enterprise AI lifecycle in Open Source



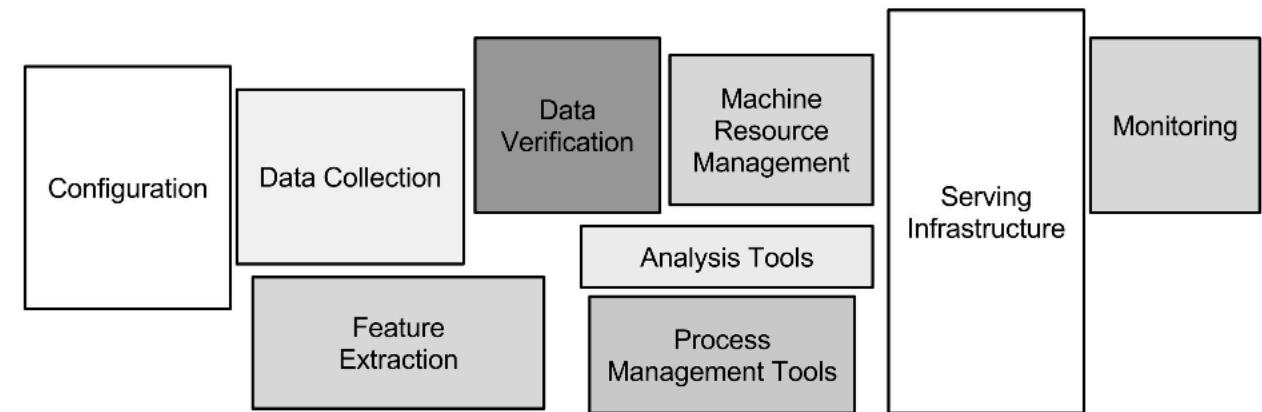
CODAIT
codait.org

The Machine Learning Workflow

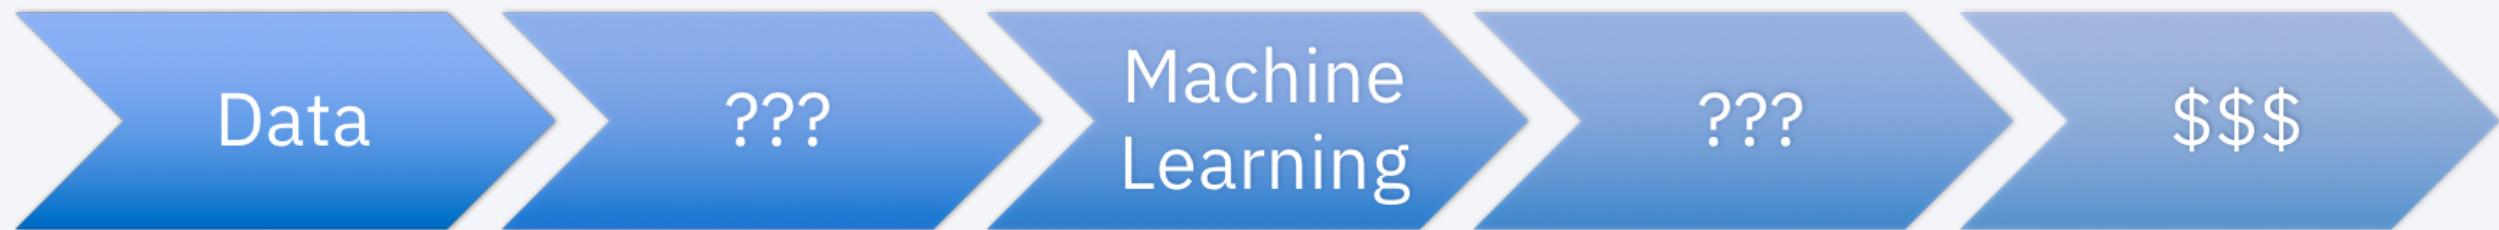


Perception

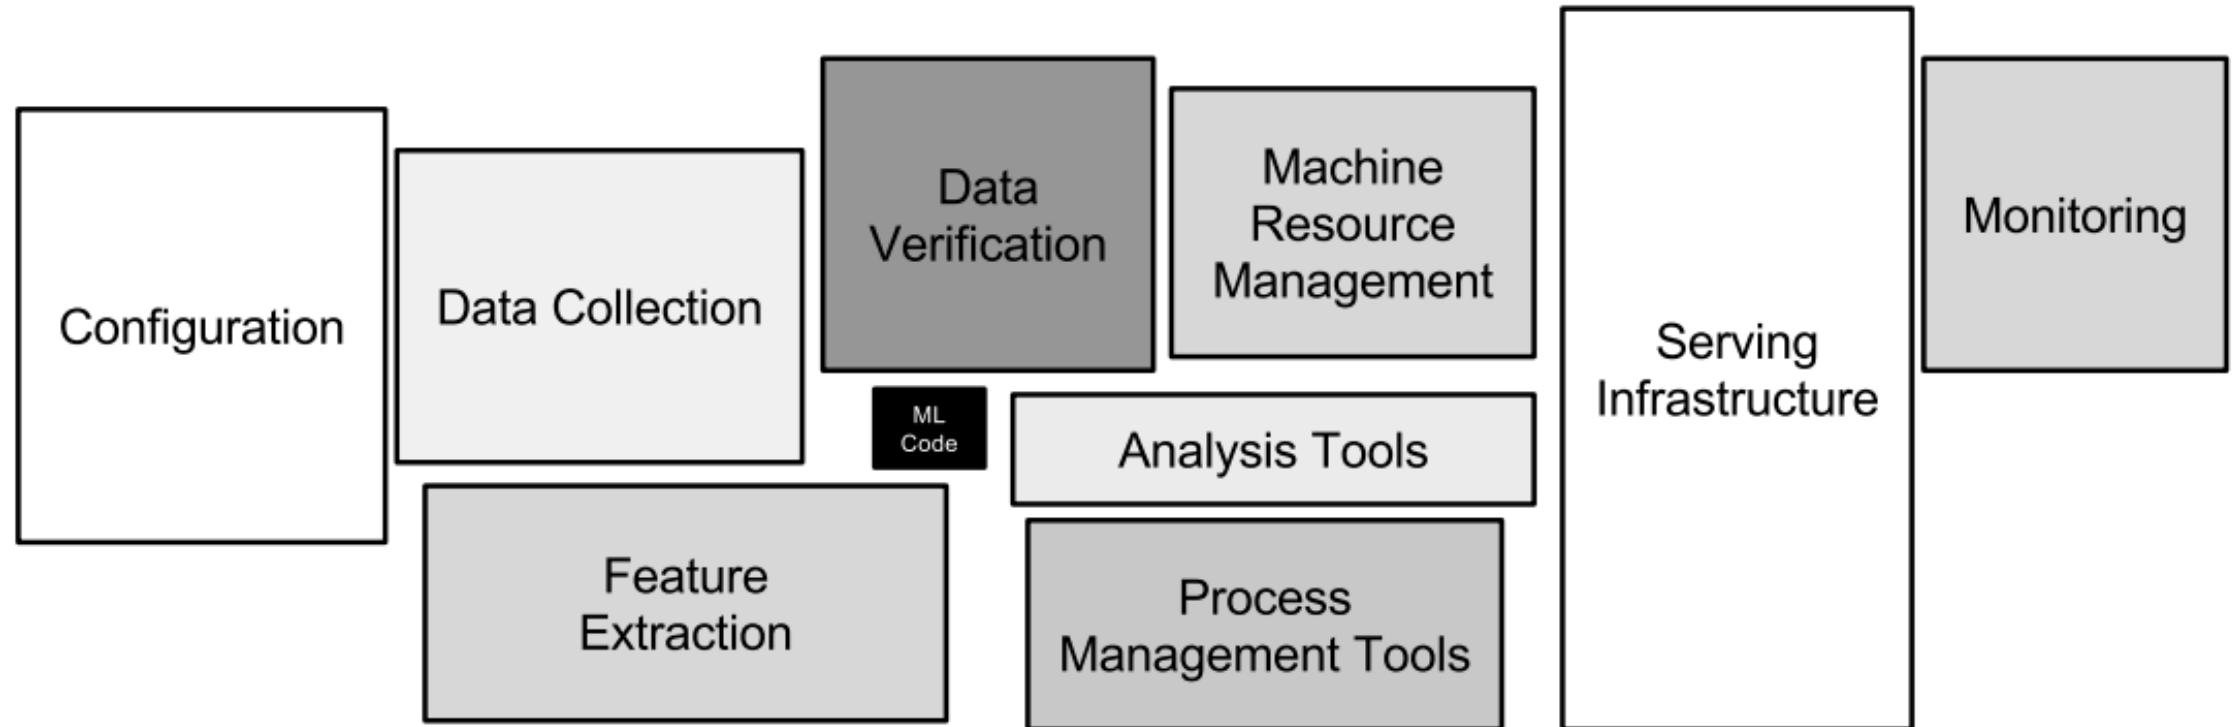
ML
Code



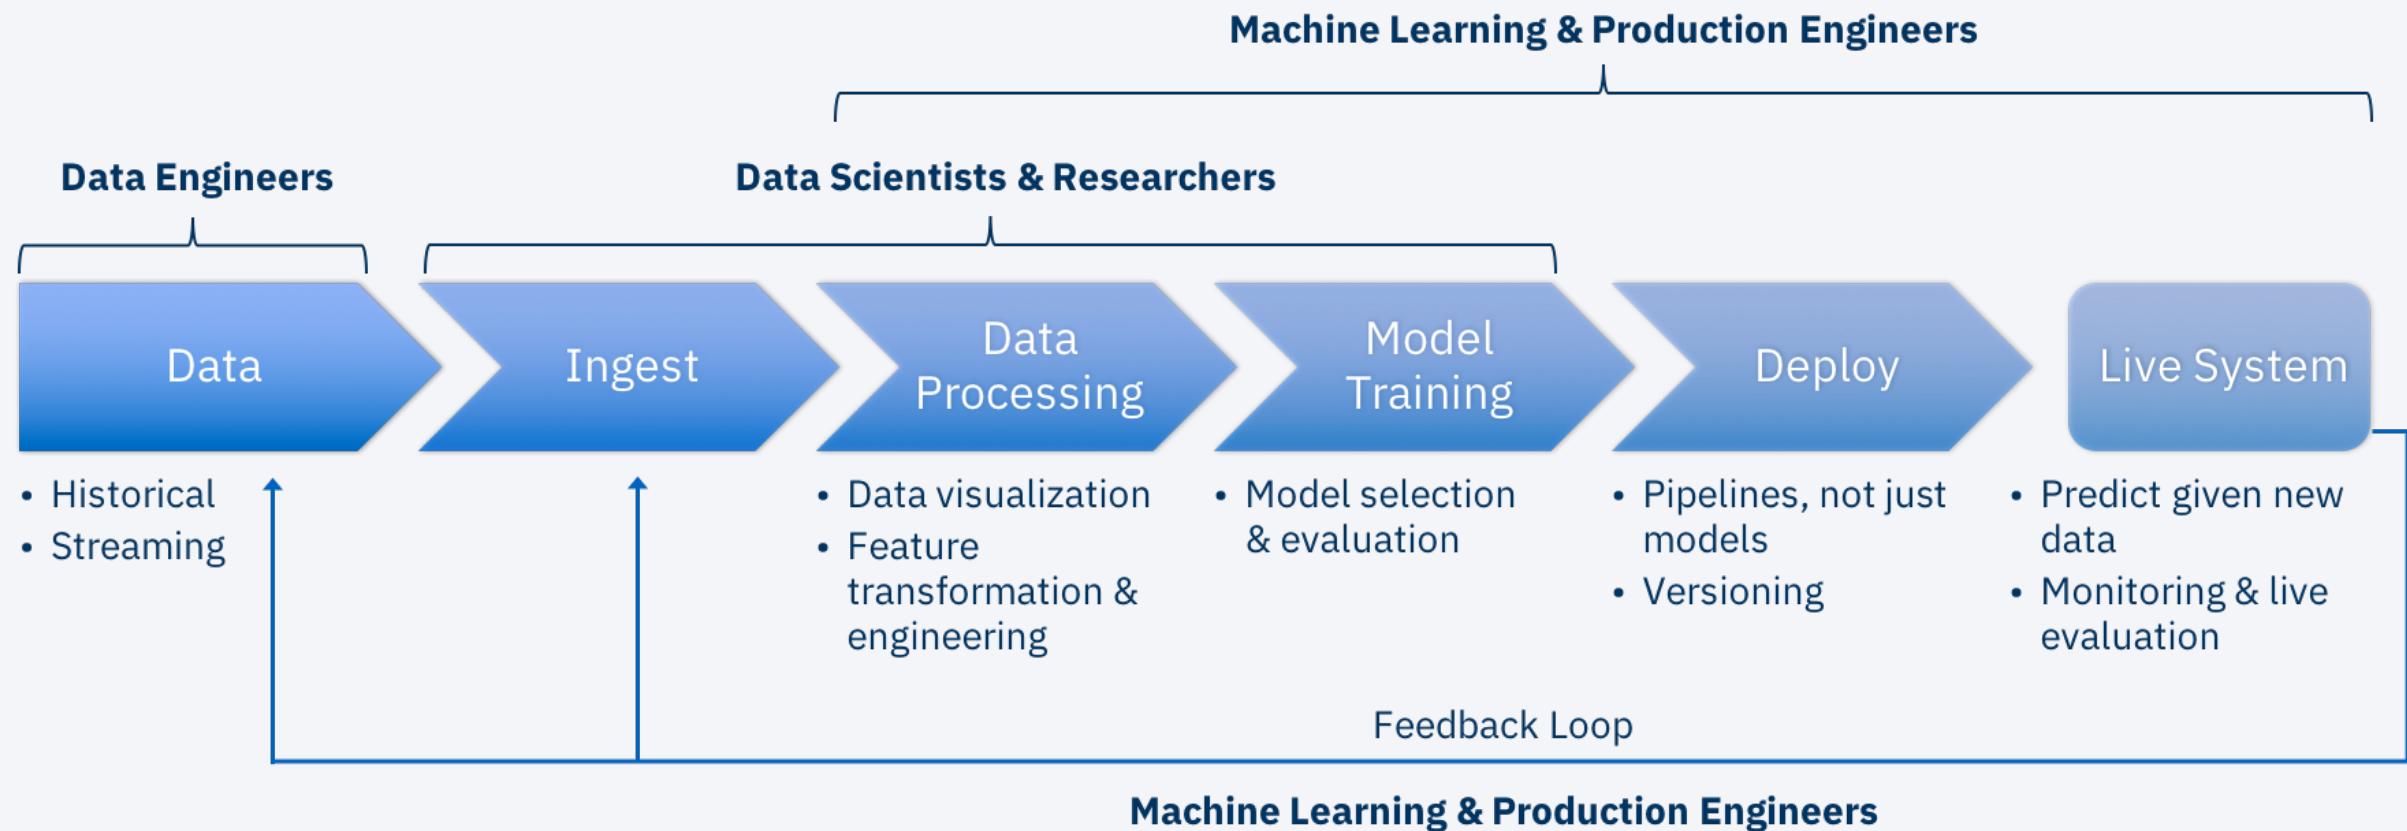
Perception



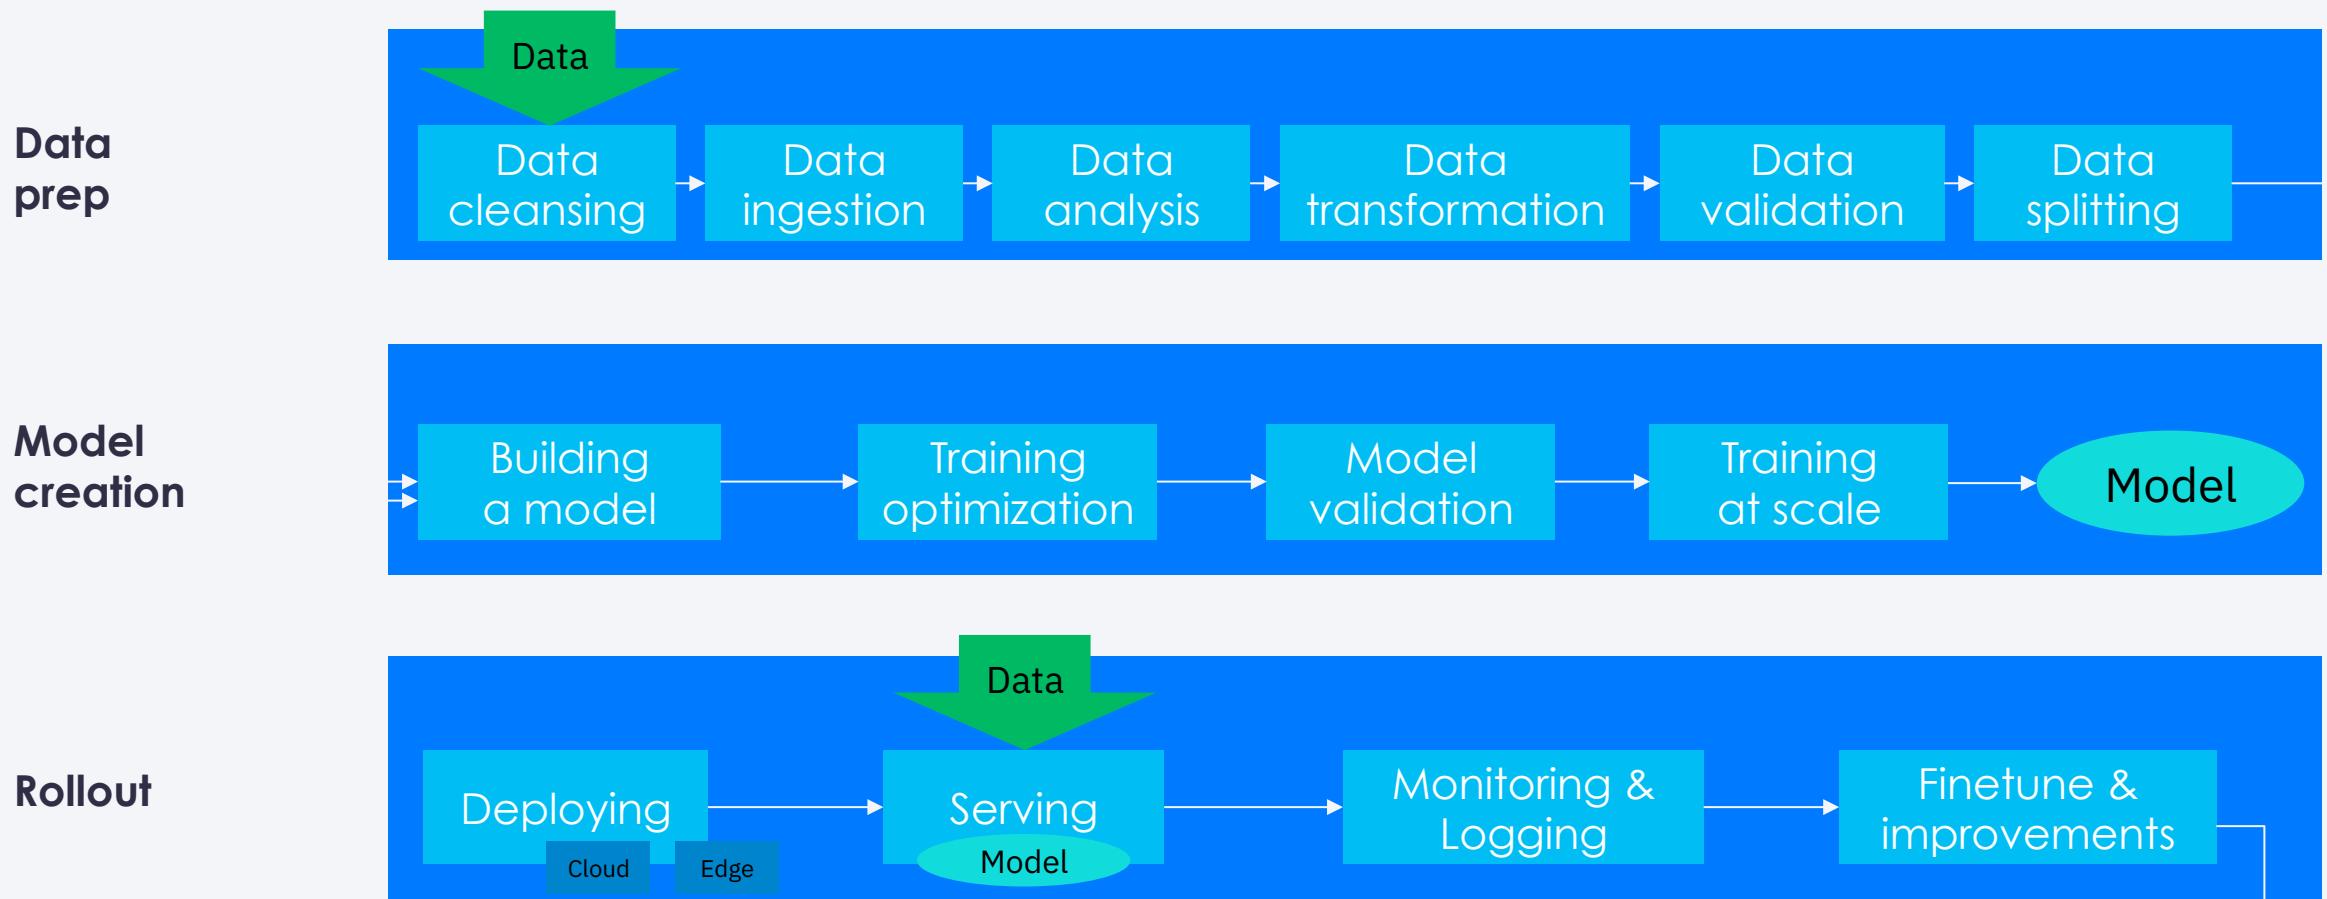
In reality...ML Code is tiny part in this overall platform



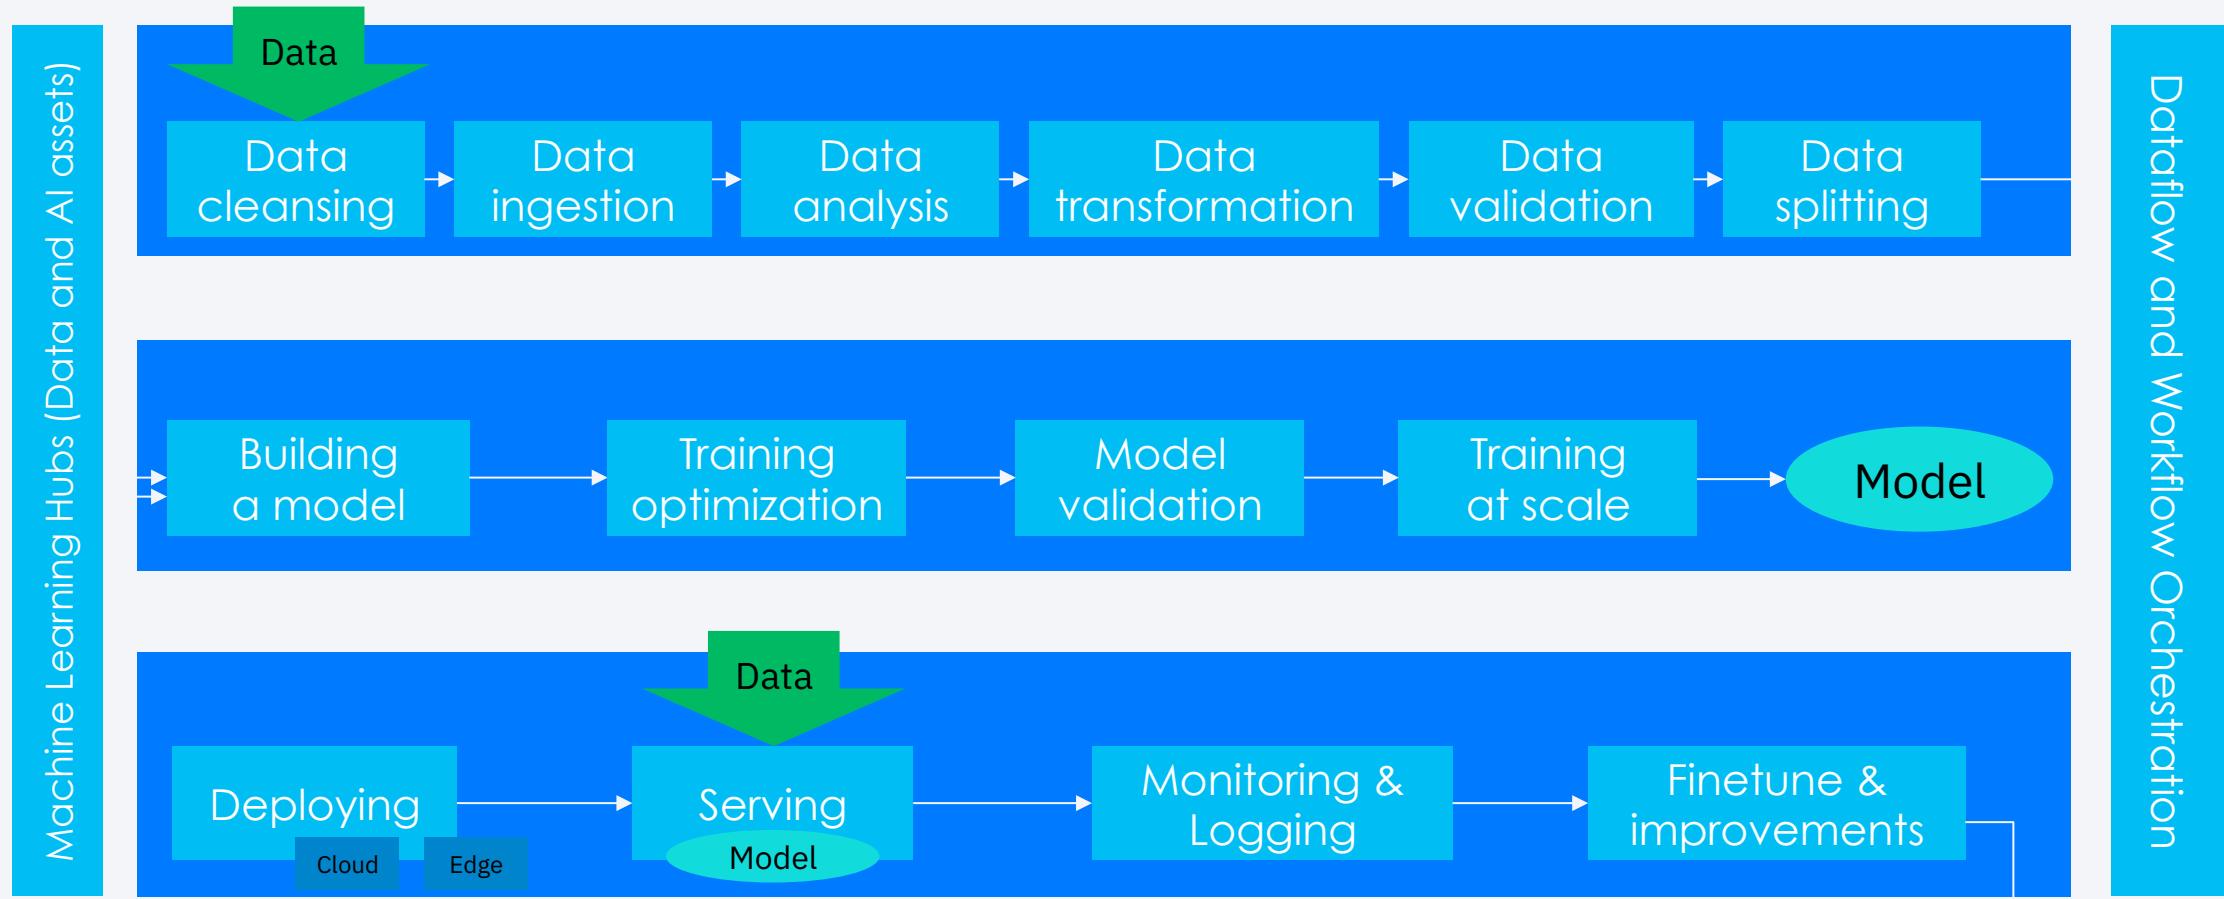
And the ML workflow spans teams ...



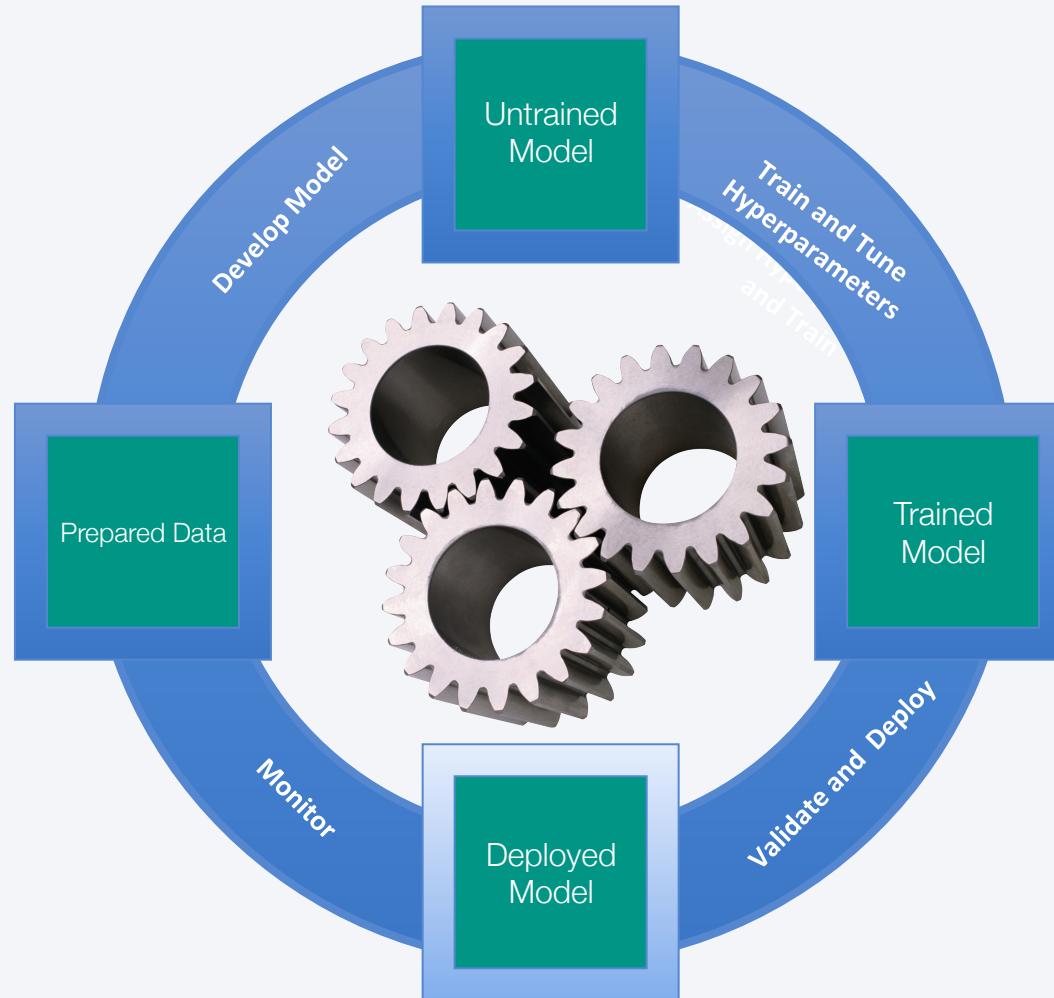
...and is much more complex...



And needs a workflow orchestrator...



Let's go through different phases of AI Lifecycle through our use cases...
For ML Lifecycle, starting with a quality data set is a must....



IBM Data Asset eXchange (DAX)

- Curated free and open datasets under open data licenses
- Standardized dataset formats and metadata
- Ready for use in enterprise AI applications

Data Asset eXchange
ibm.biz/data-asset-exchange



The screenshot shows the homepage of the IBM Data Asset eXchange (DAX). The top features a dark banner with the text "Data Asset eXchange" and "Explore useful and relevant data sets for enterprise data science". Below the banner are two buttons: "Learn More" and "Model Asset eXchange". The main content area displays six dataset cards arranged in a grid:

- Groningen Meaning Bank - Modified**: CDLA-Sharing | IOB Format. A subset of the GMB dataset, consisting of documents verified to be in the public domain. [View dataset](#). Tags: Natural Language Processing, Named Entity Recognition, Part of Speech Tagging.
- Contracts Proposition Bank**: CDLA-Sharing | CoNLL-U. Text from approximately 1000 english compliance sentences obtained from IBM's publicly available contracts, annotated with a layer of "universal" semantic role labels. [View dataset](#). Tags: Natural Language Processing, Language Modeling, Contracts.
- NOAA Weather Data - JFK Airport**: CDLA-Sharing | CSV. Local climatological data originally collected by JFK airport. [View dataset](#). Tags: Time Series Prediction.
- Nutch**: CDLA-Sharing | CSV | JSON. This dataset consists of raw and processed execution logs generated from two versions of Nutch, an open source web crawler application. [View dataset](#).
- Finance Proposition Bank**: CDLA-Sharing | CoNLL-U. Text from approximately 1000 english sentences obtained from IBM's public annual financial reports, annotated with a layer of "universal" semantic role labels. [View dataset](#).
- Forum Subjectivity**: CC BY-SA 4.0 | XML. A dataset of online discussion threads crawled from Ubuntu Forums, with associated subjectivity labels. [View dataset](#).

Data Asset eXchange

Explore useful and relevant data sets for enterprise data science

[Learn More >](#)

[Model Asset eXchange >](#)

CDLA-Sharing | CoNLL-U

Contracts Proposition Bank

Text from approximately 1000 english compliance sentences obtained from IBM's publicly available contracts, annotated with a layer of "universal" semantic role labels.

[View dataset »](#)

Natural Language Processing

Language Modeling Contracts

CDLA-Sharing | CSV

NOAA Weather Data - JFK Airport

Local climatological data originally collected by JFK airport.

[View dataset »](#)

Time Series Prediction

CDLA-Sharing | CSV | H.264

Double Pendulum Chaotic

Videos of the chaotic motion of a double pendulum, in addition to the positions of the datums of the pendulum.

[View dataset »](#)

Time Series Video Prediction

CDLA-Sharing | CSV | JSON

Nutch

CDLA-Sharing | CoNLL-U

Finance Proposition Bank

CC BY-SA 4.0 | XML

Forum Classify

Data Asset eXchange

Explore useful and relevant data sets for enterprise data science.

[Learn More >](#)

[Model Asset eXchange >](#)

CDLA-Sharing | CoNLL-U

Contracts Proposition Bank

Text from approximately 1000 english compliance sentences obtained from IBM's publicly available contracts, annotated with a layer of "universal" semantic role labels.

[Get this dataset](#)

CDLA-Sharing | CoNLL-U

Contracts Proposition Bank

Text from approximately 1000 english compliance sentences obtained from IBM's publicly available contracts, annotated with a layer of "universal" semantic role labels.

[View dataset >](#)

Natural Language Processing

Language Modeling

Contracts

CDLA-Sharing | CSV | JSON

Nutch



By [Sanjana Sahayaraj](#), [Yunyao Li](#), [Huaiyu Zhu](#), [Marina Danilevsky](#), [Poornima Chozhiyath Raman](#), [Ramiya Venkatachalam](#) | Published July 16, 2019

Natural Language Processing Language Modeling Contracts

CONTENTS

Overview

Dataset Metadata

Example Records

Citation

RESOURCES

Data Asset eXchange (DAX)
Explore useful and relevant data sets for enterprise data science.

Model Asset eXchange (MAX)
A place for developers to find and use free and open source deep learning models.

Center for Open-Source Data & AI Technologies (CODAIT)
Improving the Enterprise AI Lifecycle in Open Source.

Overview

ConProp version 1.0 was developed by researchers at IBM Almaden Research Center, San Jose, CA, USA. It consists of proposition bank-style annotations from approximately 1000 english compliance sentences obtained from IBM's publicly available contracts. These sentences were extracted from contract sections such as Business Partner descriptions, Agreement Terms / Structure, Intellectual Property Protection, Limitation of Liability, Warranty Terms, General Principles of Relationship, Terms of Agreement Termination, Withdrawal of Service, Third Party Claims, Charges, Service Level Agreement and many more.

To suit the need in compliance domain, a list of 60 predicates specific to the domain were

Data Asset eXchange

Explore useful and relevant data sets for enterprise data science.

[Learn More >](#)

[Model Asset eXchange >](#)

CDLA-Sharing | CoNLL-U

Contracts Proposition Bank

Text from approximately 1000 english compliance sentences obtained from IBM's publicly available contracts, annotated with a layer of "universal" semantic role labels.

[Get this dataset](#)

IBM Cloud Object Storage

CDLA-Sharing | CoNLL-U

Contracts Proposition Bank

Text from approximately 1000 english compliance sentences obtained from IBM's publicly available contracts, annotated with a layer of "universal" semantic role labels.

[View dataset >](#)

Natural Language Processing

Language Modeling

Contracts

CDLA-Sharing | CSV | JSON

Nutch



By [Sanjana Sahayaraj](#), [Yunyao Li](#), [Huaiyu Zhu](#), [Marina Danilevsky](#), [Poornima Chozhiyath Raman](#), [Ramiya Venkatachalam](#) | Published July 16, 2019

Natural Language Processing

Language Modeling

Contracts

CONTENTS

Overview

Dataset Metadata

Example Records

Citation

RESOURCES

Data Asset eXchange (DAX)
Explore useful and relevant data sets for enterprise data science.

Model Asset eXchange (MAX)
A place for developers to find and use free and open source deep learning models.

Center for Open-Source Data & AI Technologies (CODAIT)
Improving the Enterprise AI Lifecycle in Open Source.

Overview

ConProp version 1.0 was developed by researchers at IBM Almaden Research Center, San Jose, CA, USA. It consists of proposition bank-style annotations from approximately 1000 english compliance sentences obtained from IBM's publicly available contracts. These sentences were extracted from contract sections such as Business Partner descriptions, Agreement Terms / Structure, Intellectual Property Protection, Limitation of Liability, Warranty Terms, General Principles of Relationship, Terms of Agreement Termination, Withdrawal of Service, Third Party Claims, Charges, Service Level Agreement and many more.

To suit the need in compliance domain, a list of 60 predicates specific to the domain were

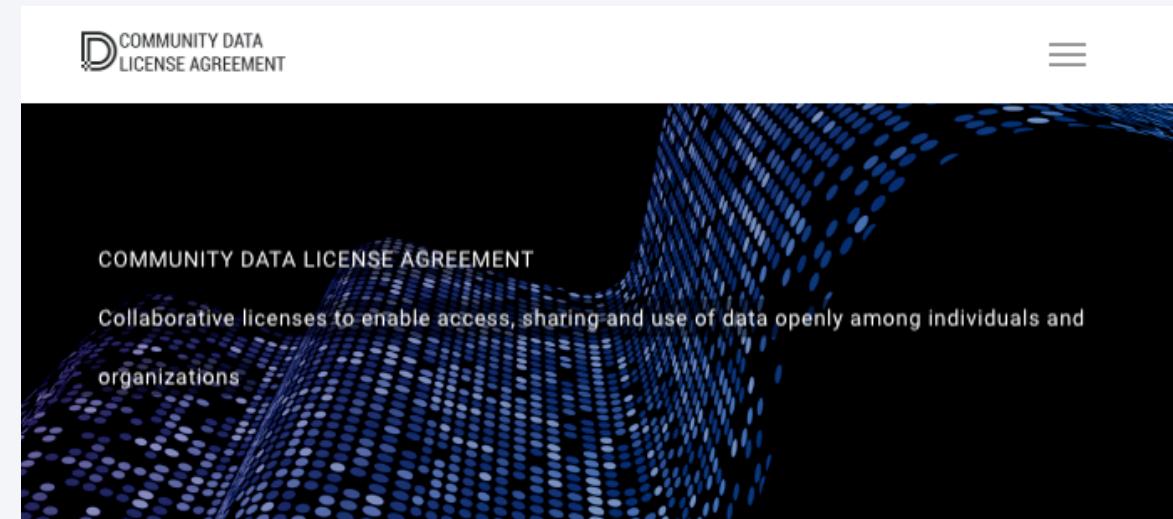
The Community Data License Agreement, driven by Linux Foundation

<http://cdla.io>

Two licenses written specifically for AI data

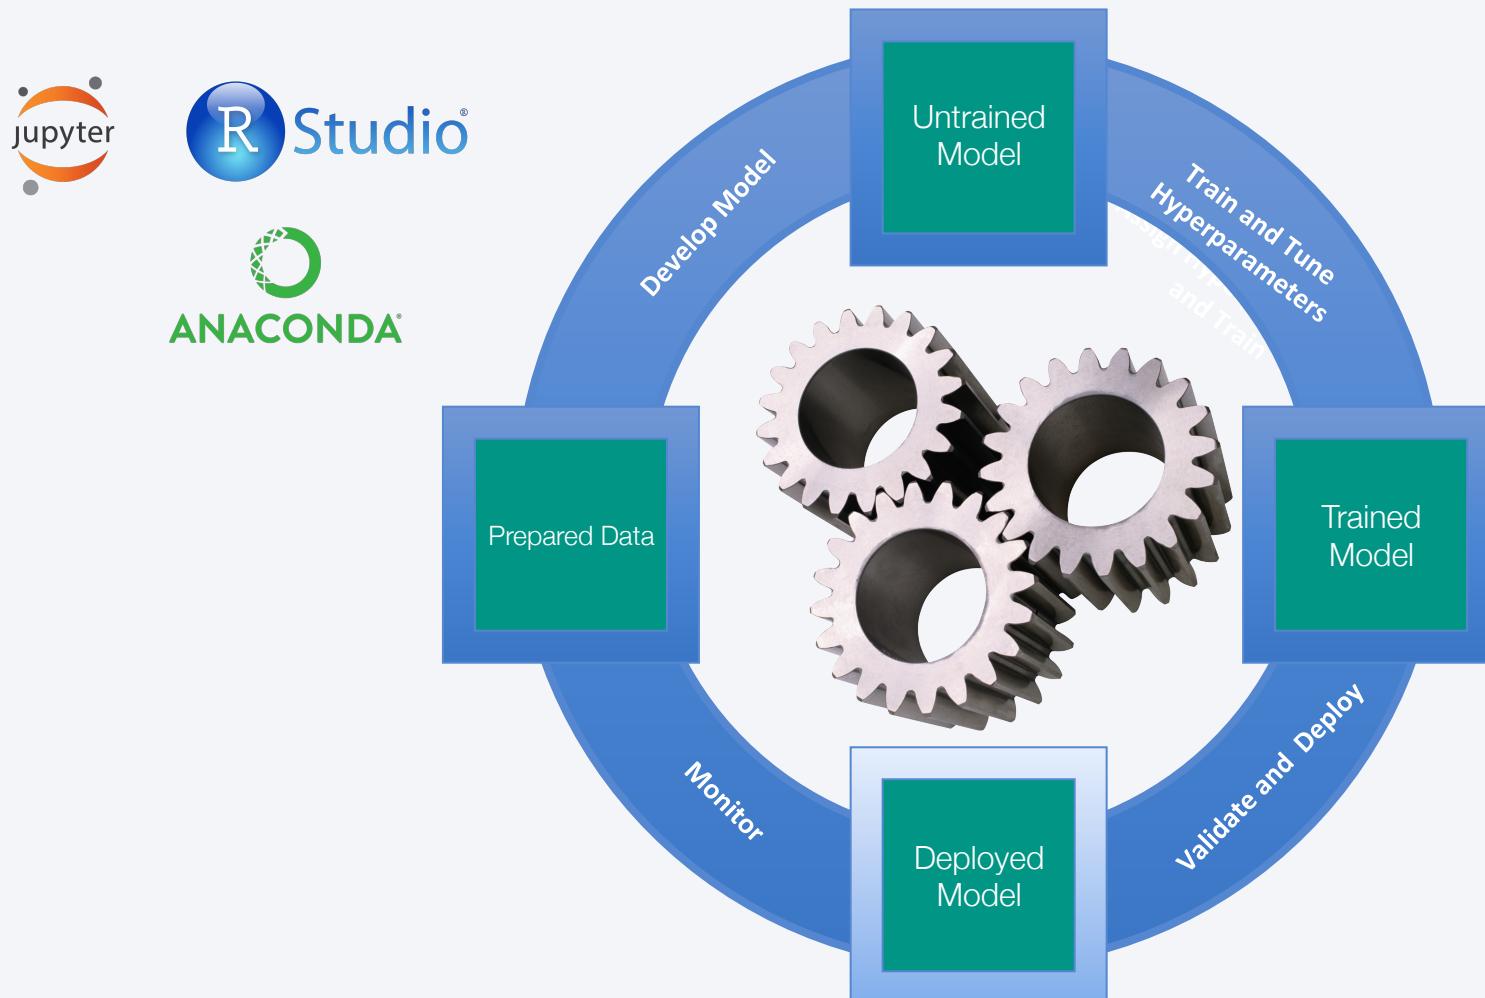
- CDLA-Sharing: “Copyleft” license analogous the GPL
- CDLA-Permissive: Similar to BSD license

Both licenses distinguish clearly between *use* (analysis, modeling) and *modification* of the data set.

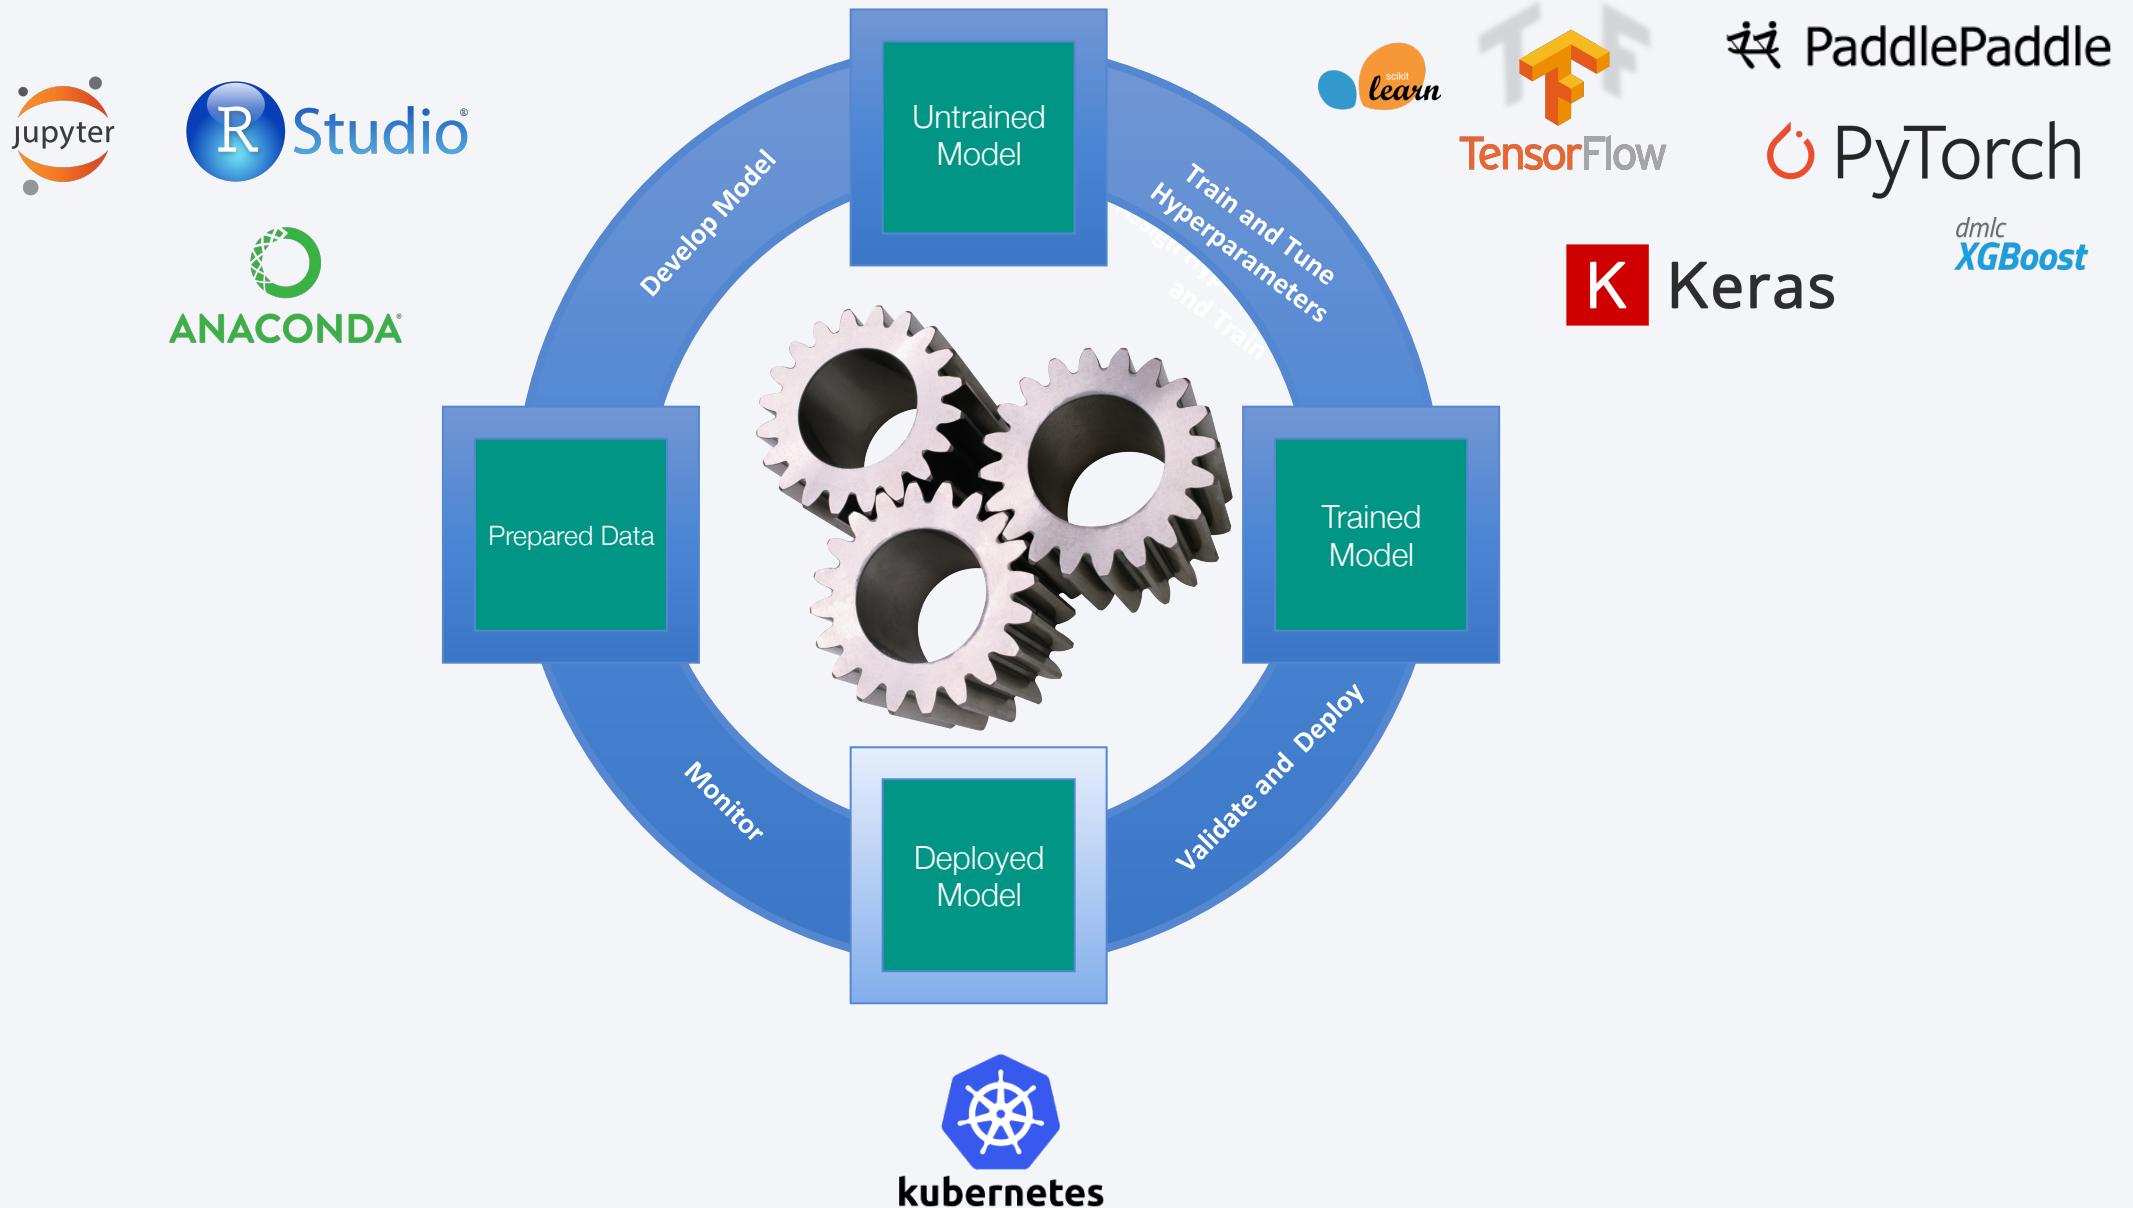


Open source software communities have shown the power of open collaboration building some of the world's most important software assets together. There are communities also looking to collaboratively build datasets that can be shared and developed in a very similar model to software. For example, machine learning and AI systems require vast amounts of training data. Governments are looking for ways to establish public-private sharing of data.

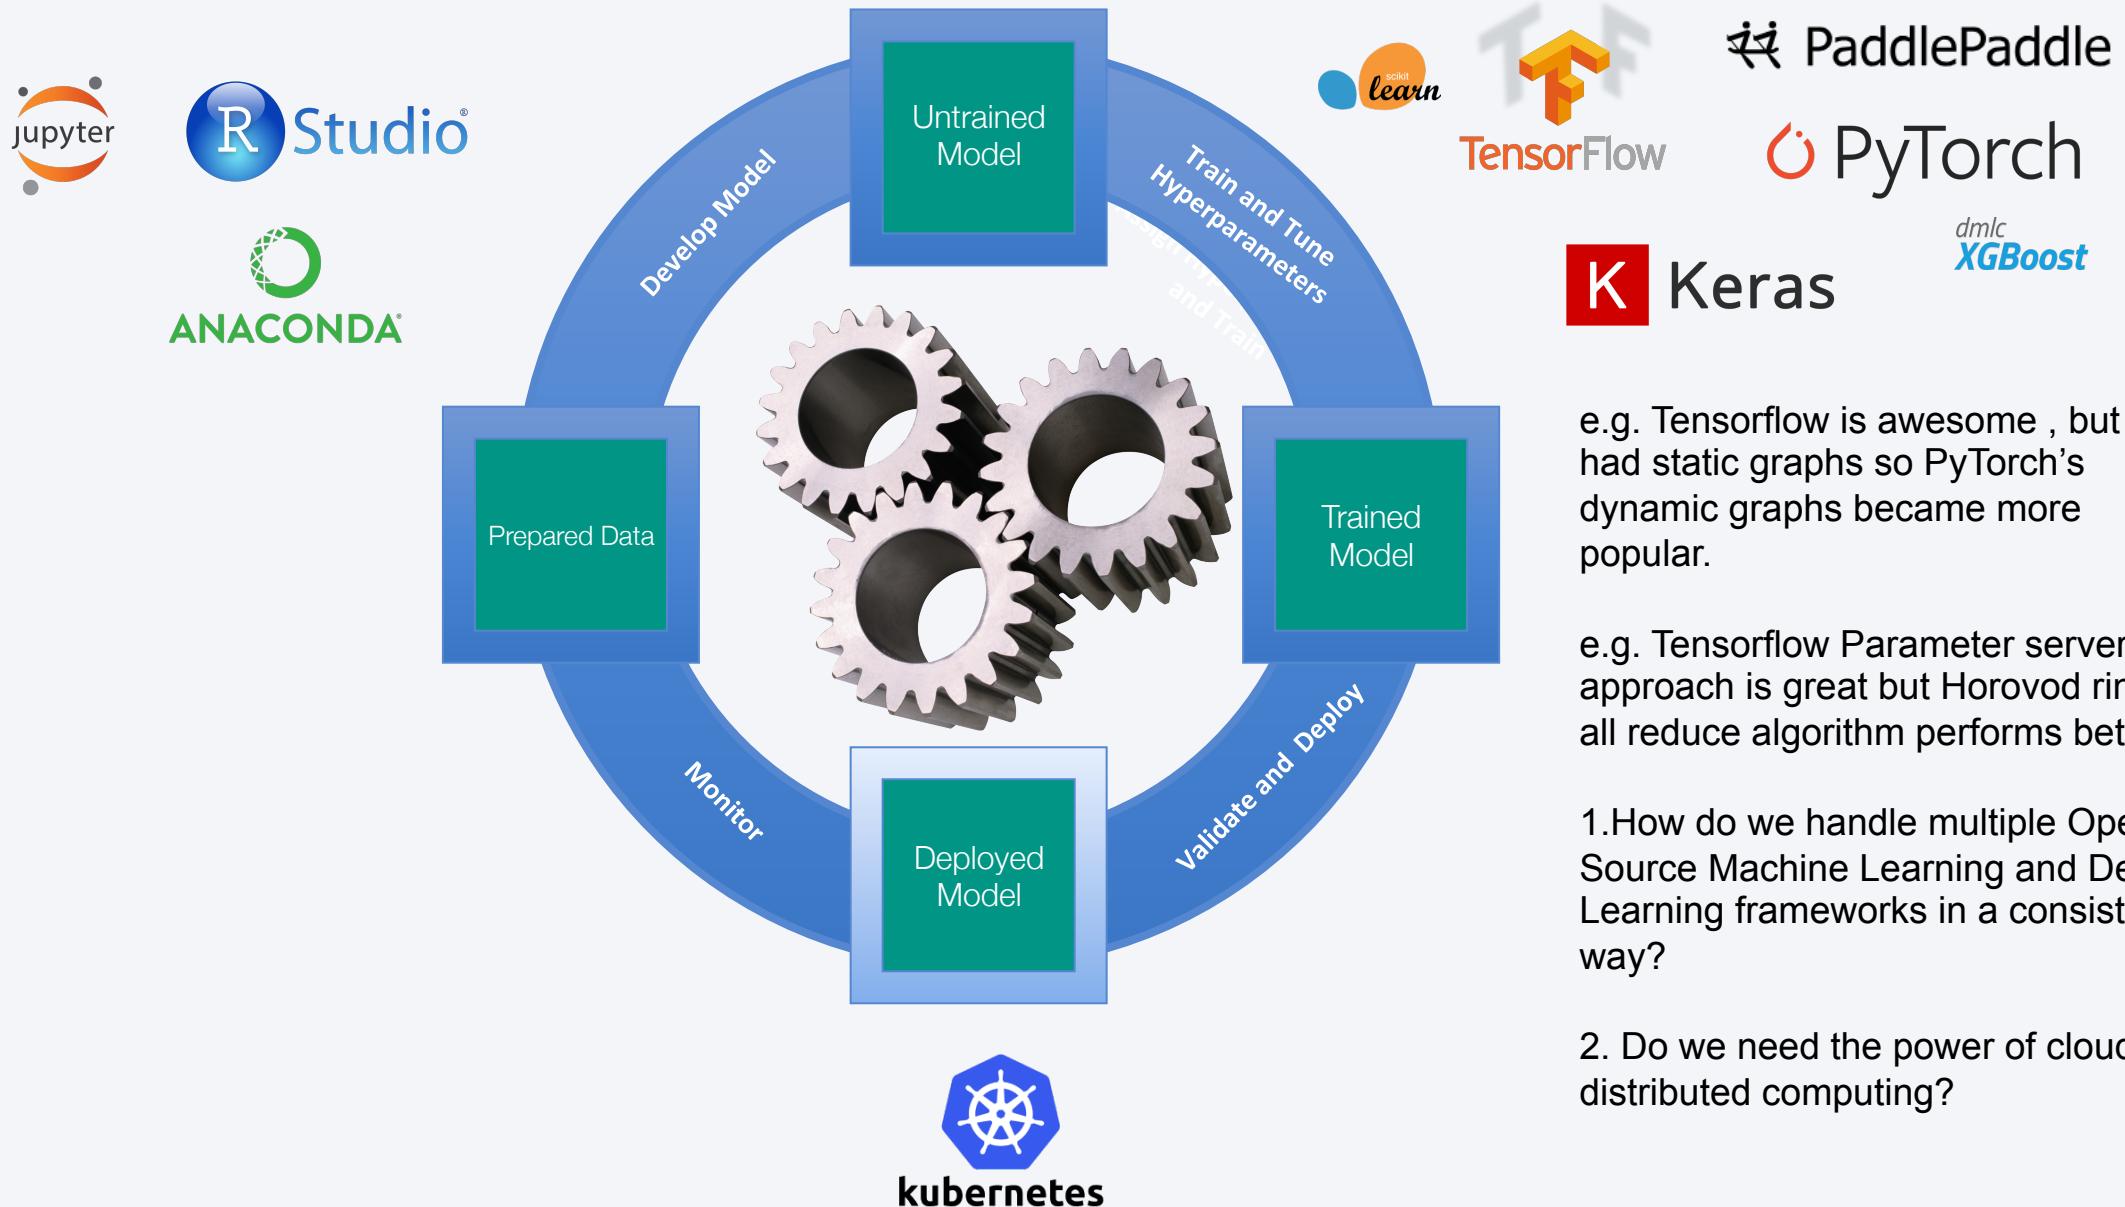
Many tools available to build initial models



Many tools to train machine learning and deep learning models



We need a multi framework ML platform. And it should be cloud native... but why?



e.g. Tensorflow is awesome , but had static graphs so PyTorch's dynamic graphs became more popular.

e.g. Tensorflow Parameter server approach is great but Horovod ring-all reduce algorithm performs better?

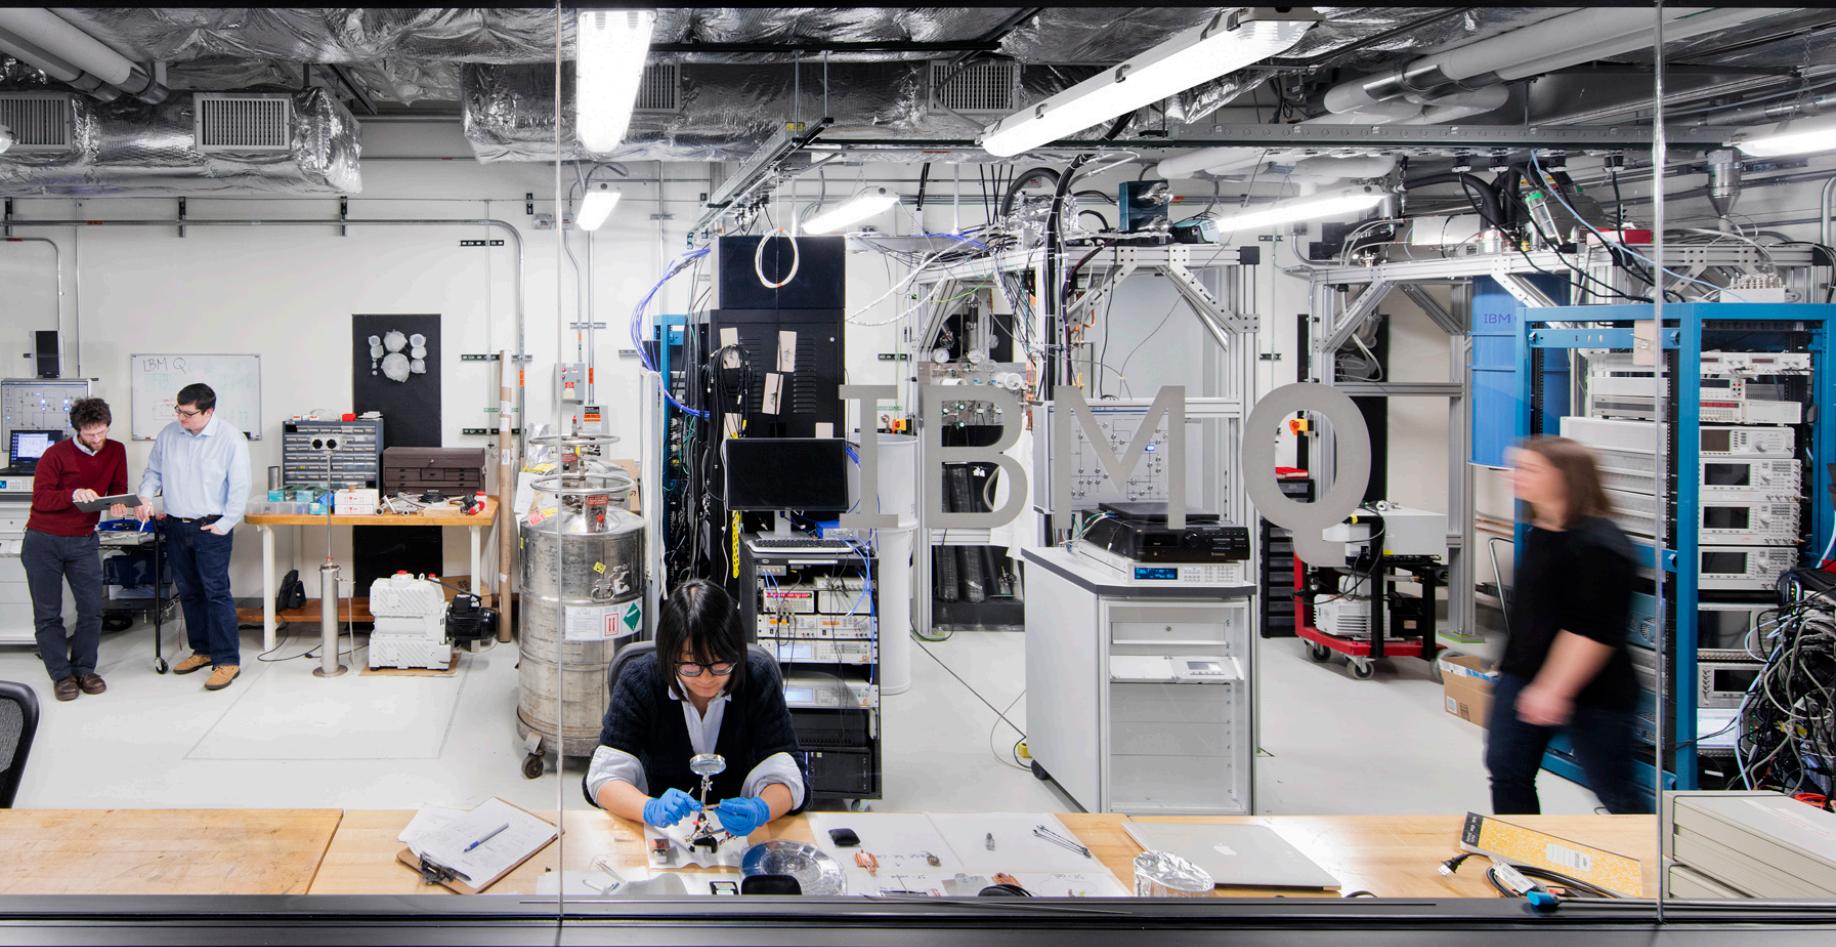
1.How do we handle multiple Open Source Machine Learning and Deep Learning frameworks in a consistent way?

2. Do we need the power of cloud for distributed computing?

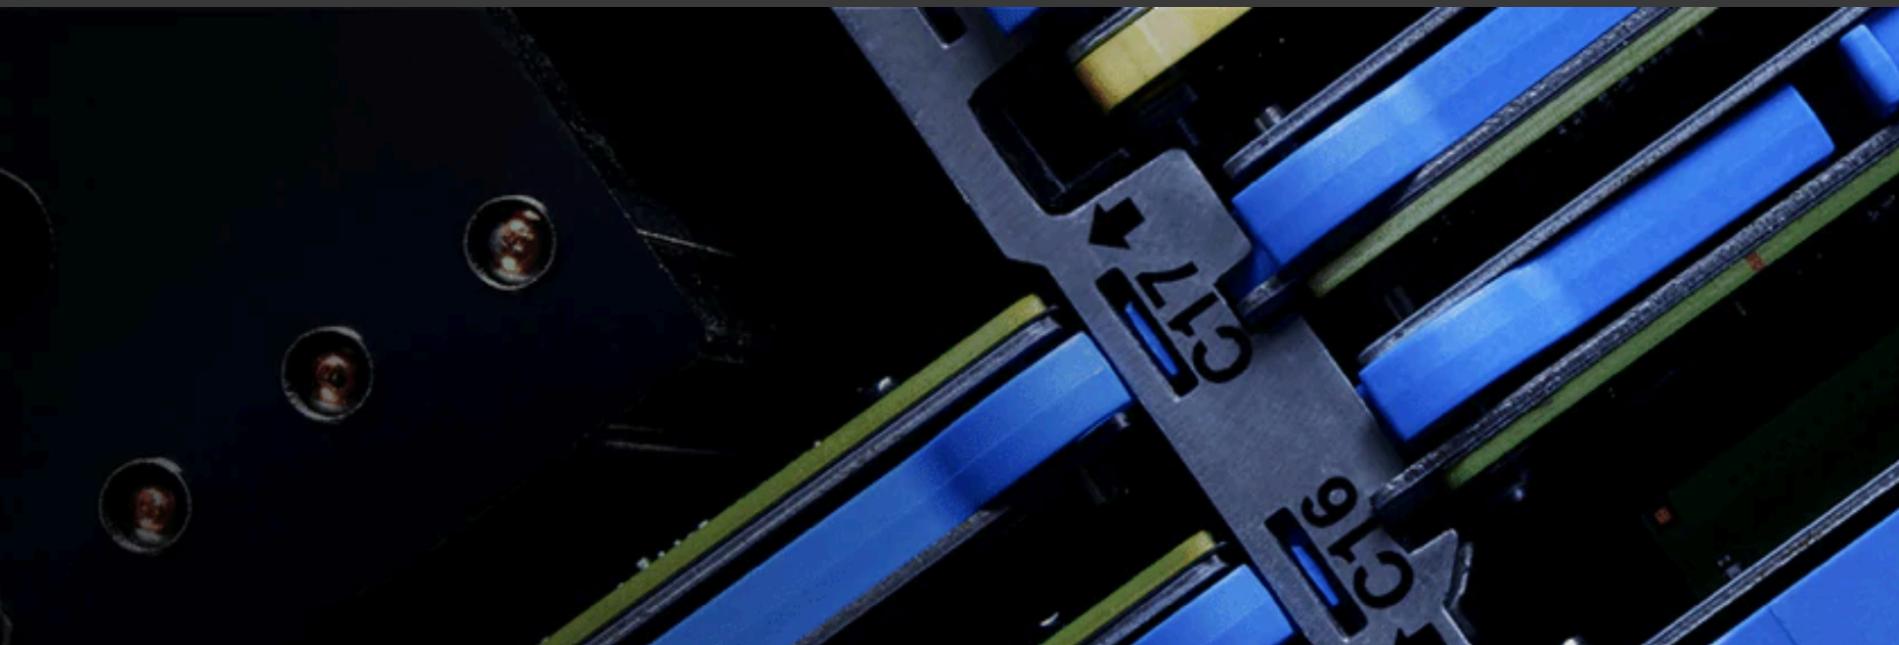
ML requires
the strength
of HPC &
GPUs



Ability to
scale ML
workloads on
demand



Ability to utilize various technologies and achieve high performance computing.



1. Model/Data Parallelism
2. MPI/NCCL

NCCL (pronounced "Nickel") is a stand-alone library of standard collective communication routines for GPUs, implementing all-reduce, all-gather, reduce, broadcast, and reduce-scatter. It has been optimized to achieve high bandwidth on platforms using PCIe, NVLink, NVswitch, as well as networking using InfiniBand Verbs or TCP/IP sockets.

NCCL supports an arbitrary number of GPUs installed in a single node or across multiple nodes, and can be used in either single- or multi-process (e.g., MPI) applications.

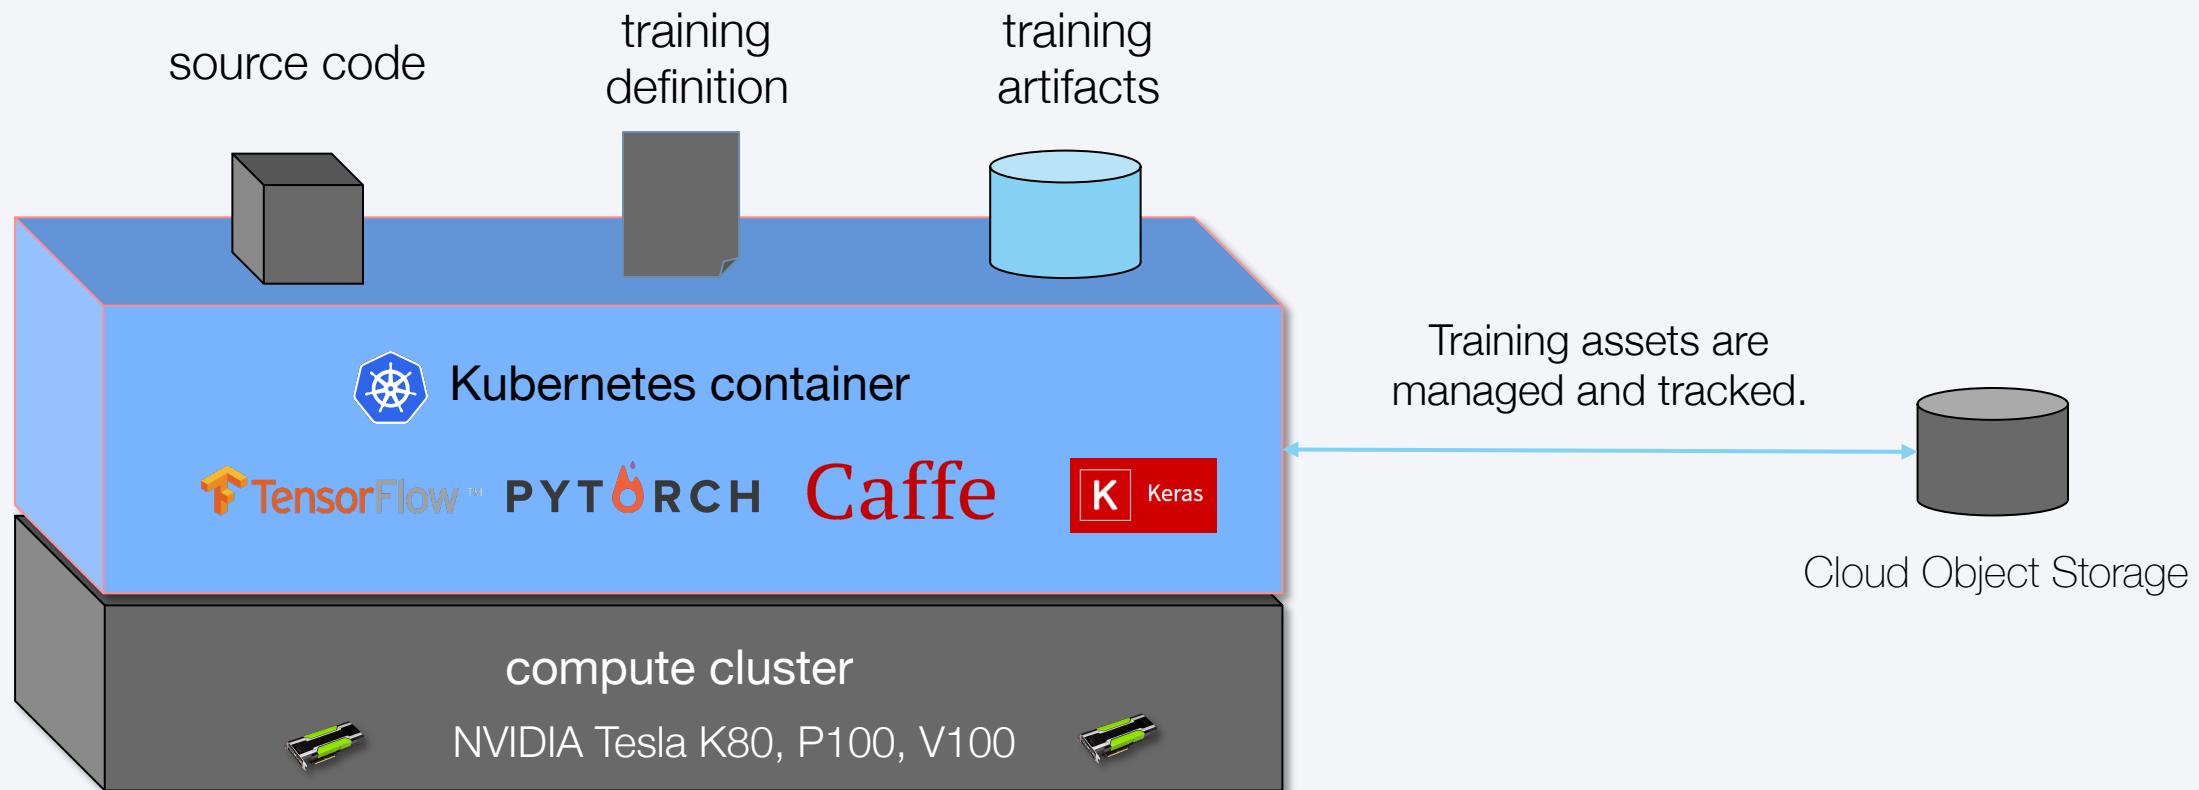


To scale we need to go
Cloud native for AI....we need to
go to Kubernetes

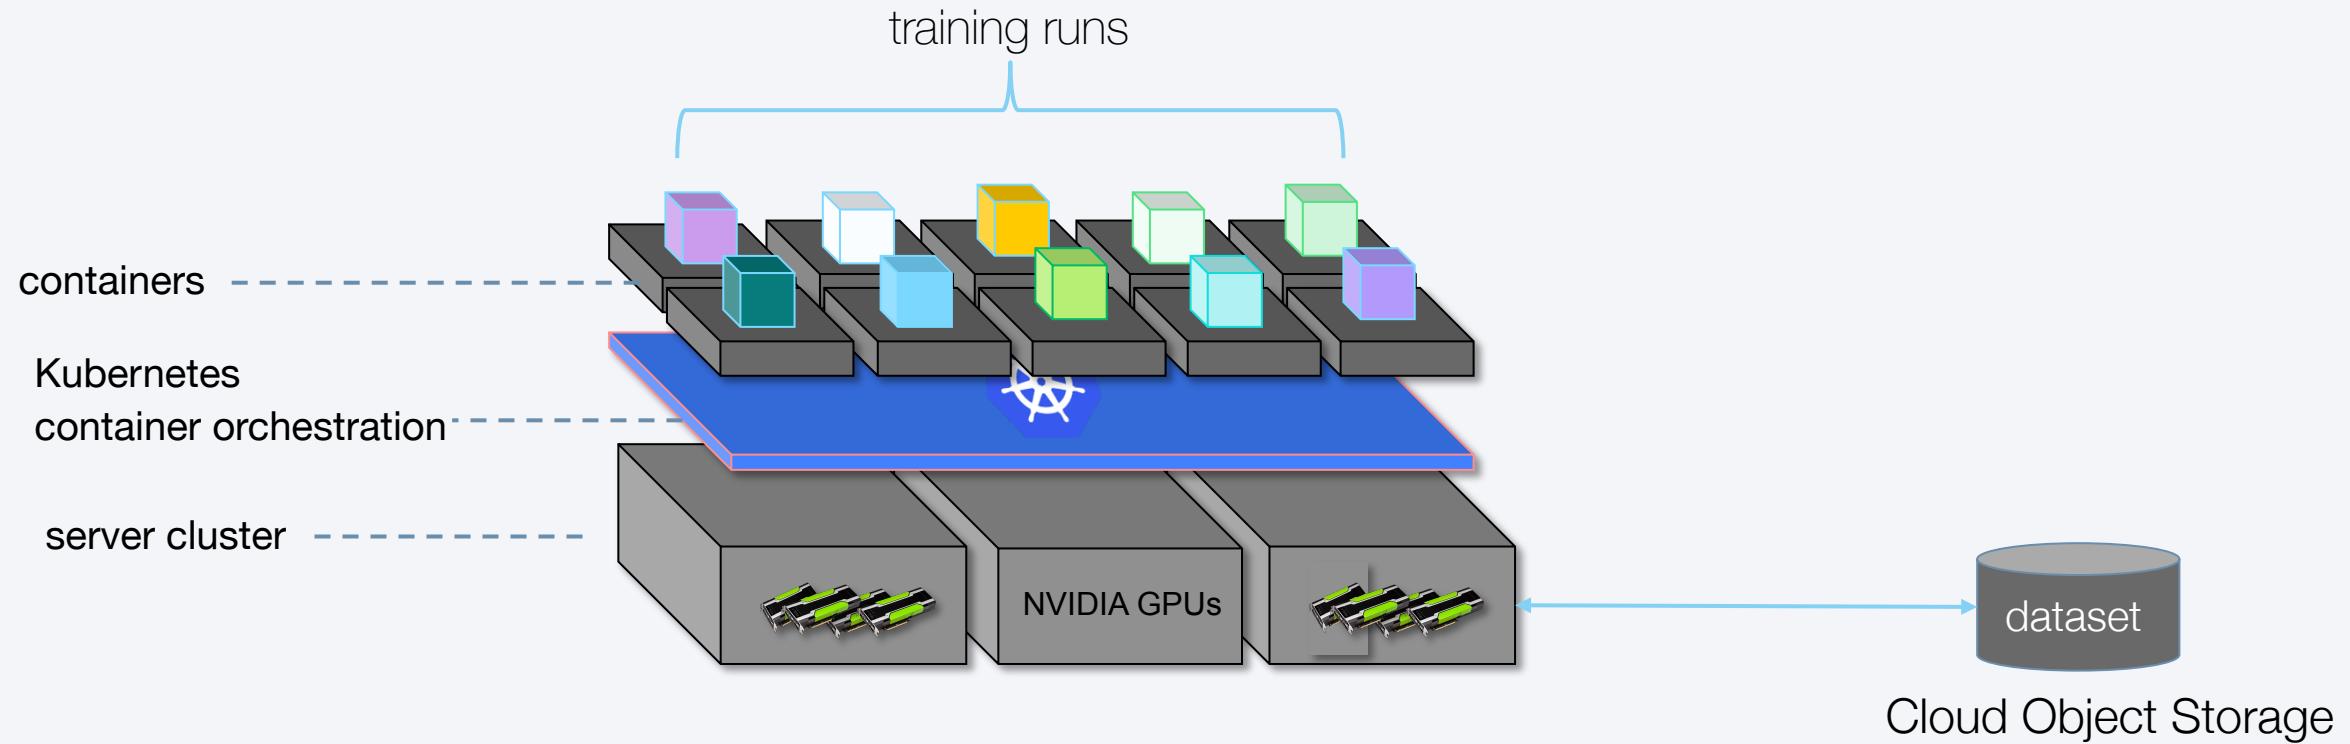
Microservices
Containers
DevOps
automation

Access to elastic compute leveraging Kubernetes

Auto-allocation means infrastructure is used only when needed



Model training distributed across containers



So we need Kubernetes, but Kubernetes is not the end game



Kelsey Hightower

@kelseyhightower

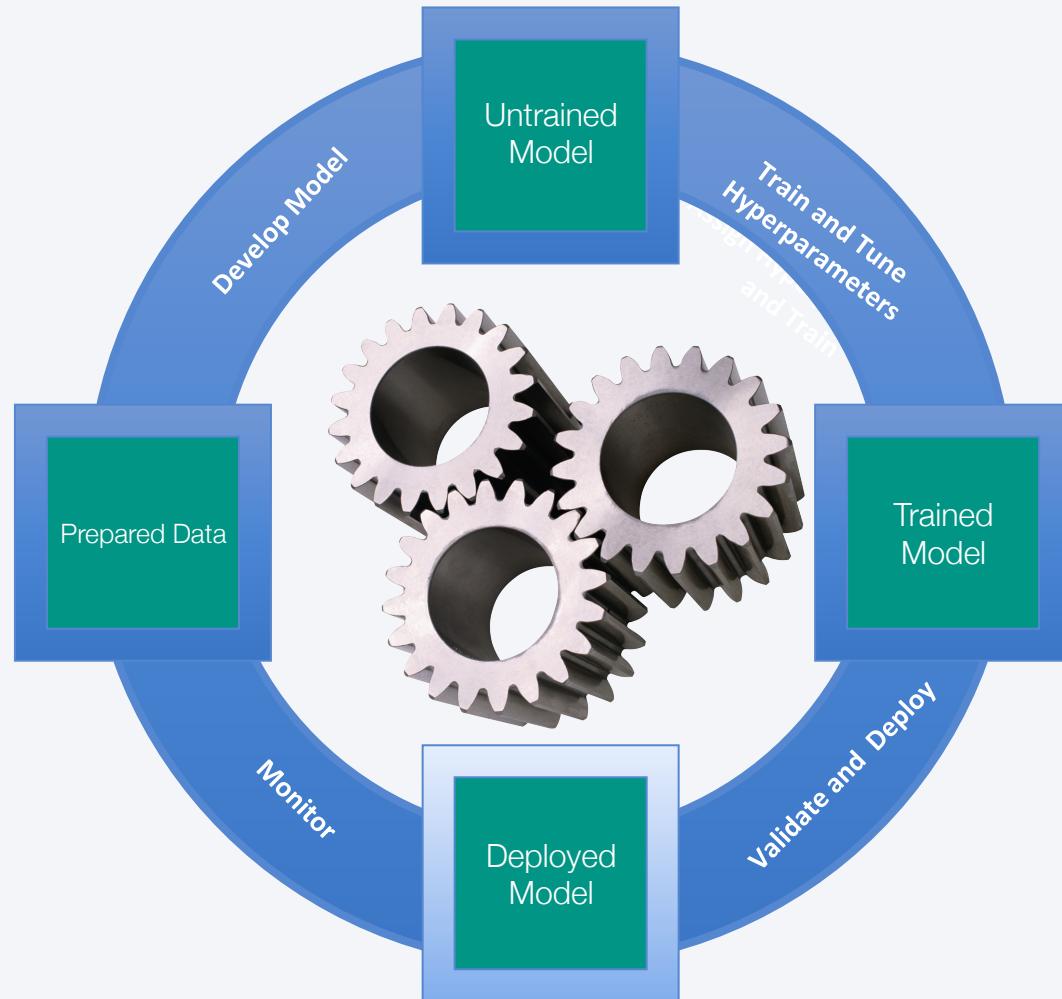


Kubernetes is a platform for building platforms. It's a better place to start; not the endgame.

556 1:04 PM - Nov 27, 2017



We need a platform. Enter Kubeflow



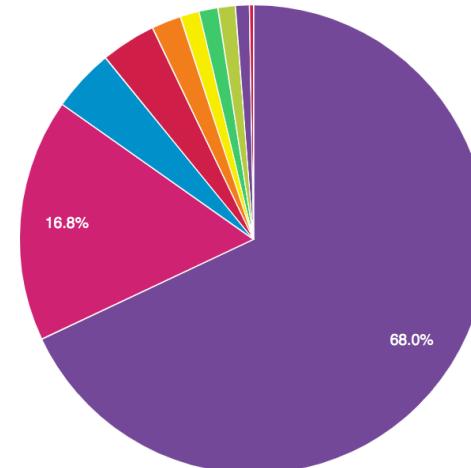
Kubeflow

<https://github.com/kubeflow>

- End to end ML Platform on Kubernetes. Focused on multiple aspects of Model Lifecycle
- Originated at Google, and has grown to have a large community of developers
- Google, IBM, Cisco, RedHat, Intel, Microsoft and others contributing
- IBM is the 2nd largest contributor in terms of overall commits. IBM maintainers (committers/reviewers) in Katib (HPO+Training), Kubeflow Serving, Manifests, Pipelines etc.

Commits by Company

Show 10 entries		Search
#	Company	Commits
	*Independent	3669
1	Google	904
2	IBM	234
3	Cacloud	205
4	Alibaba	110
5	Red Hat	71
6	Intel	65
7	Huawei	52
8	Cisco Systems	15
9	VA Linux	12



Libraries and CLIs - Focus on end users

Arena

kubectl

kubectl

fairing

Systems - Combine multiple services

katib

pipelines

Model DB

kube
bench

notebooks

TFX

Low Level APIs / Services (single function)

TFJob

PyTorchJ
ob

Pipelines CR

Argo

Jupyter
CR

MPI CR

Seldon CR

Study Job

Spark Job

Developed By
Kubeflow

Developed Outside
Kubeflow

* Not all components shown

<https://ibm.biz/bootstrap>

IAM

Orchestration

Scheduling

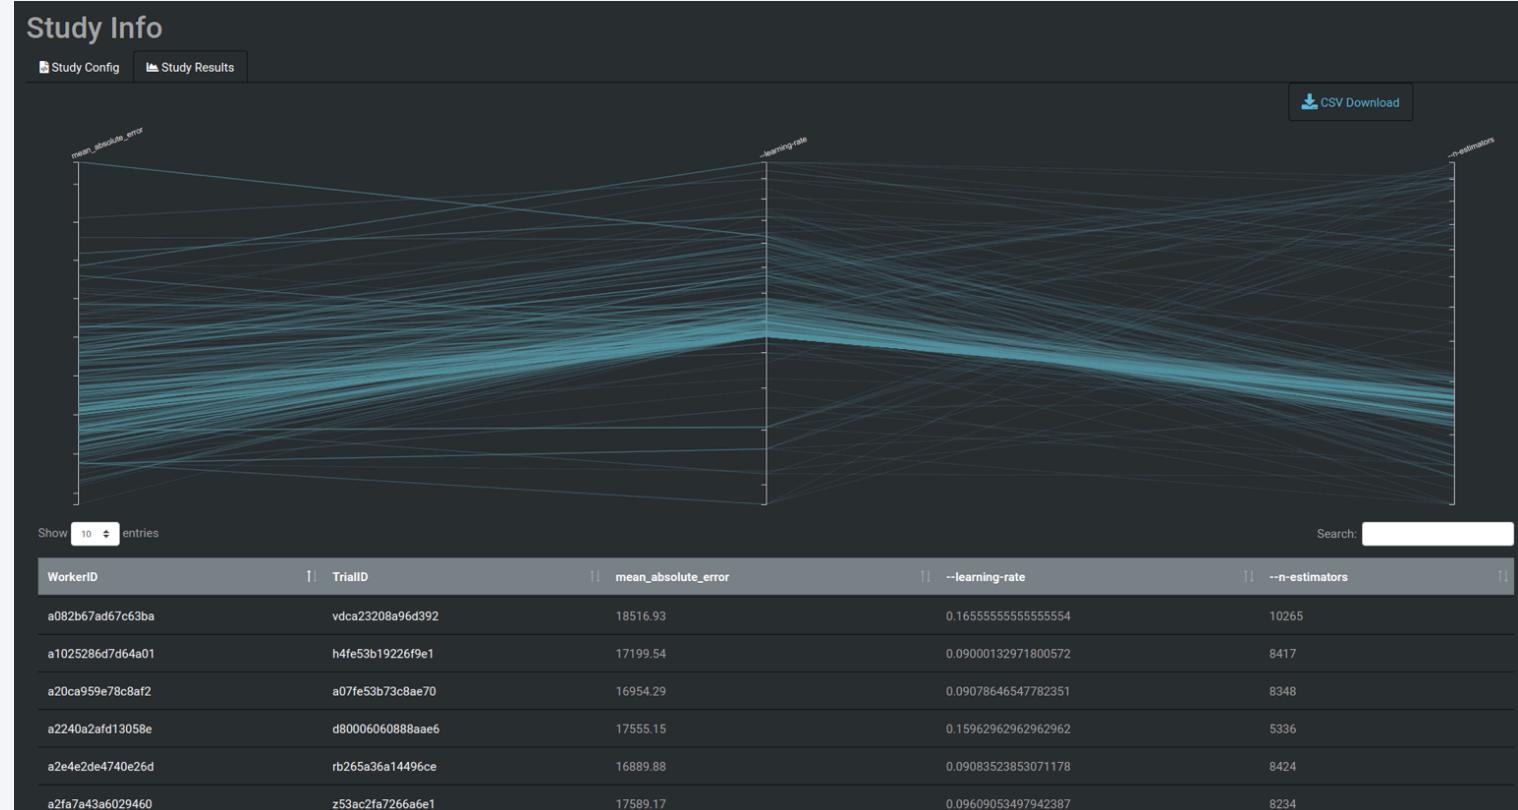
Metadata

Distributed Model Training and HPO (Katib, TFJob, PyTorch Job...)

- Addresses One of the key goals for model builder persona:

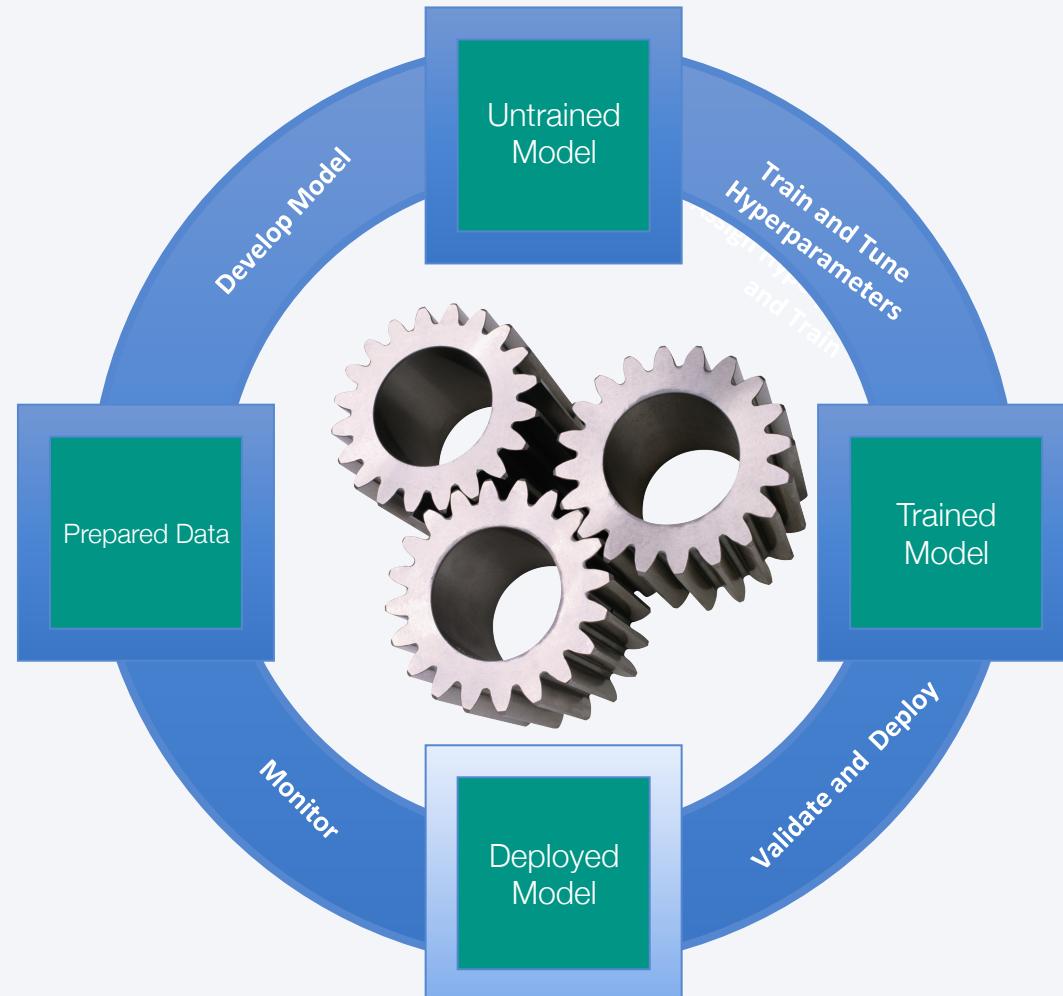
Distributed Model Training and Hyper parameter optimization for Tensorflow, PyTorch etc.

- Common problems in HP optimization
 - Overfitting
 - Wrong metrics
 - Too few hyperparameters
- Katib: a fully open source, Kubernetes-native hyperparameter tuning service
 - Inspired by Google Vizier
 - Framework agnostic
 - Extensible algorithms
 - Simple integration with other Kubeflow components
- Kubeflow also supports distributed MPI based training using Horovod



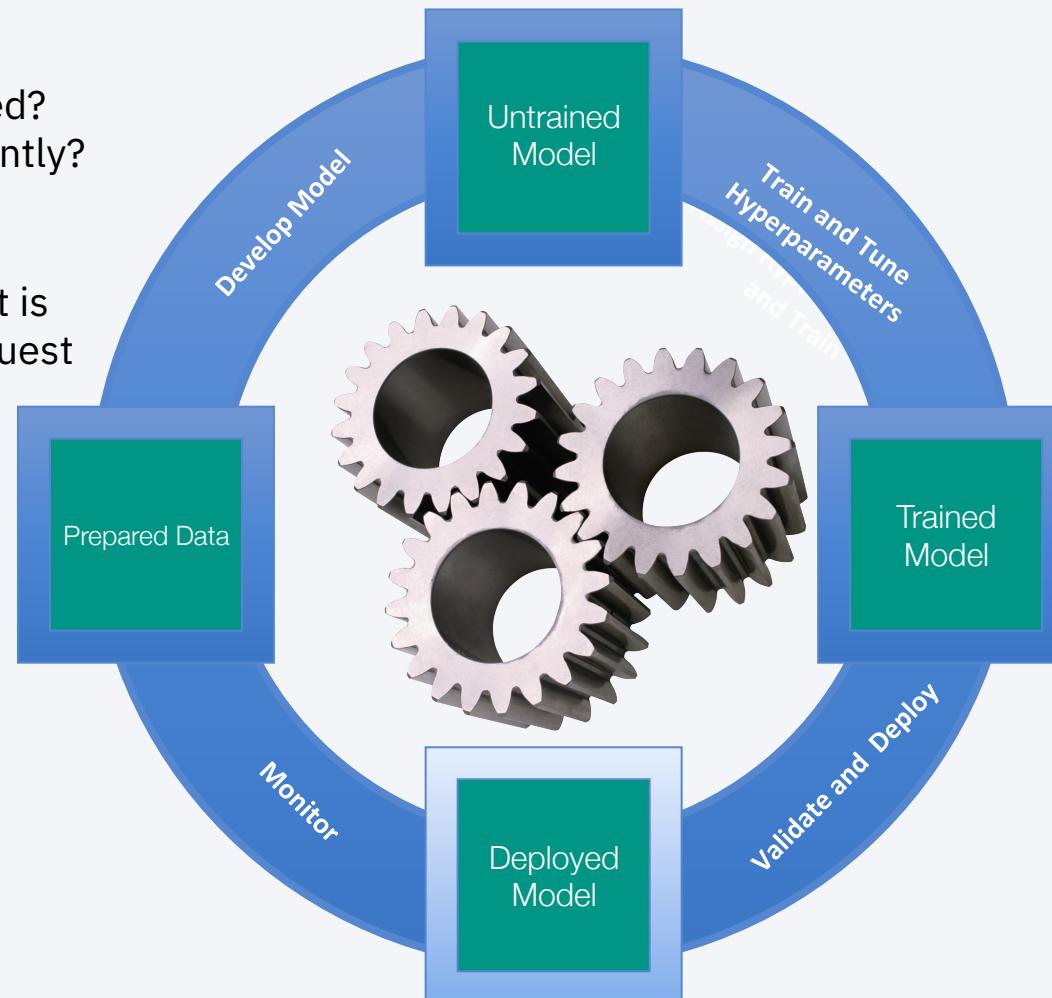
Model is ready for deployment: Just wrap in containers and push on Kube?

Anything else needed?



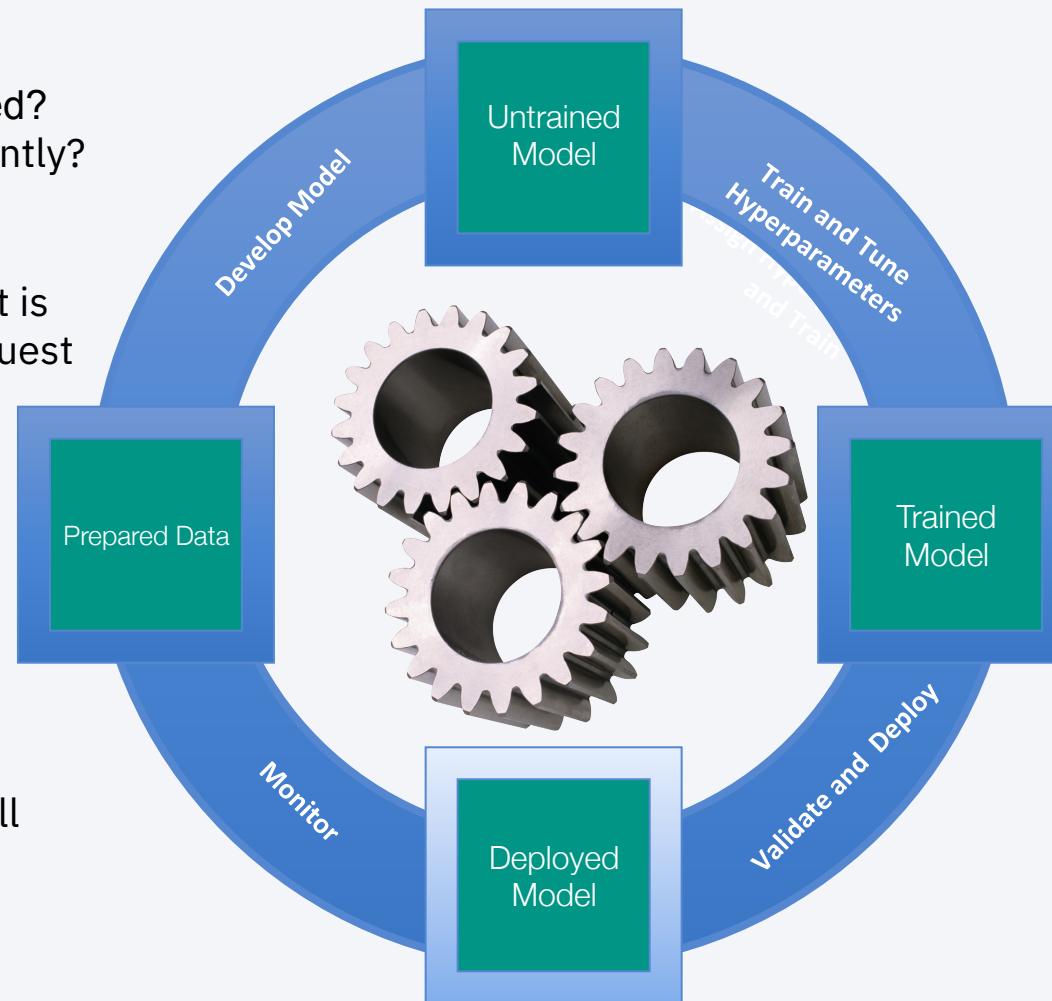
Production ML Model Serving? How hard could it be?

- ❑ Cost:
Is the model over or under scaled?
Are resources being used efficiently?
- ❑ Monitoring:
Are the endpoints healthy? What is
the performance profile and request
trace?



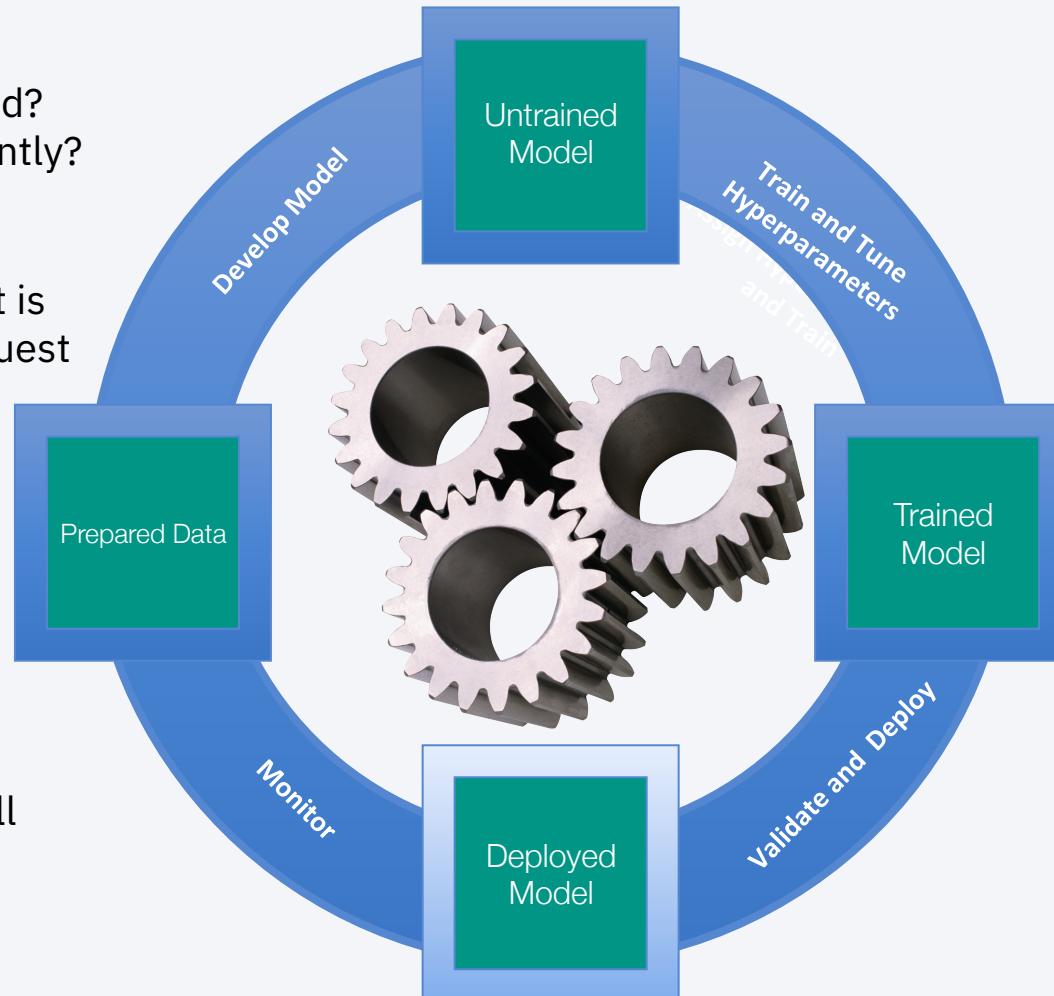
Production ML Model Serving? How hard could it be?

- Cost:
Is the model over or under scaled?
Are resources being used efficiently?
- Monitoring:
Are the endpoints healthy? What is
the performance profile and request
trace?
- Rollouts:
Is this rollout safe? How do I roll
back? Can I test a change
without swapping traffic?
- Protocol Standards:
How do I make a prediction?
GRPC? HTTP? Kafka?



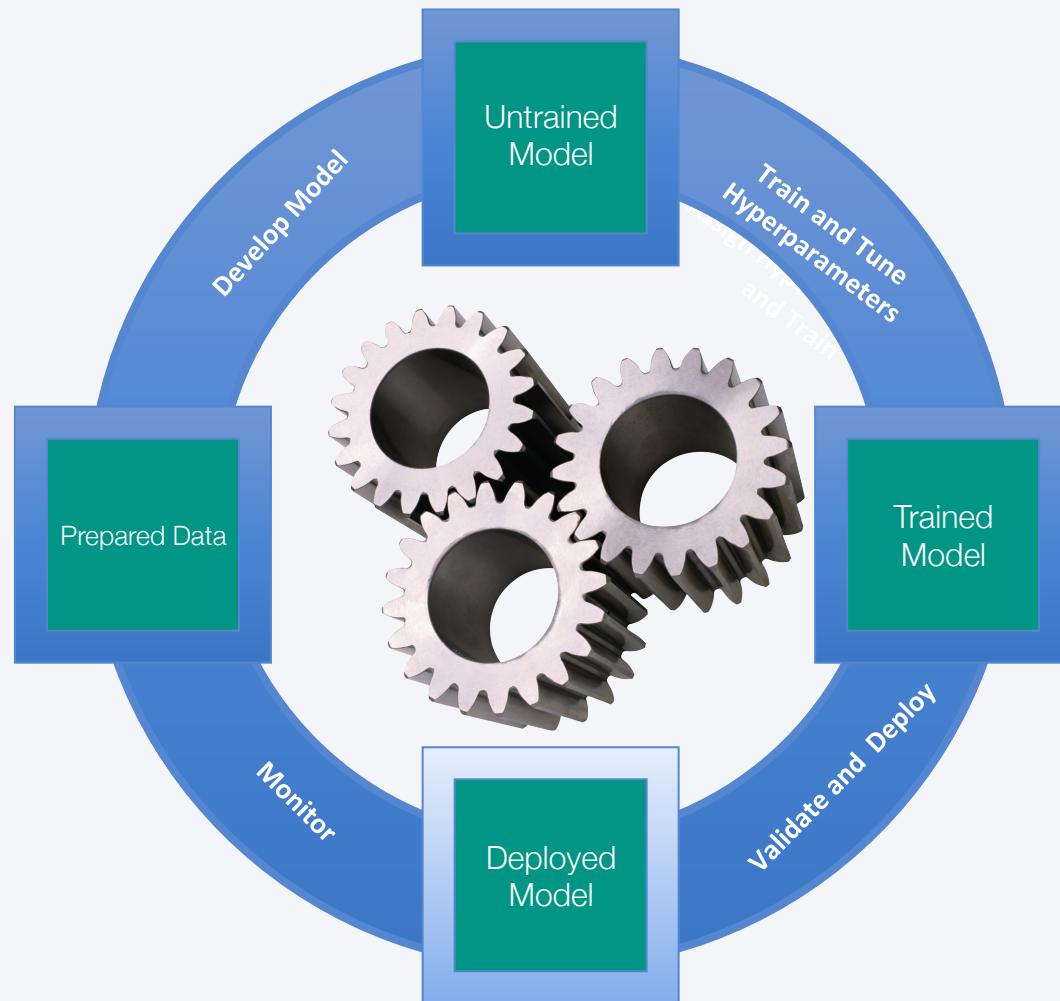
Production ML Model Serving? How hard could it be?

- Cost:**
Is the model over or under scaled?
Are resources being used efficiently?
- Monitoring:**
Are the endpoints healthy? What is the performance profile and request trace?
- Rollouts:**
Is this rollout safe? How do I roll back? Can I test a change without swapping traffic?
- Protocol Standards:**
How do I make a prediction?
GRPC? HTTP? Kafka?



- Frameworks:**
How do I serve on Tensorflow?
XGBoost? Scikit Learn? Pytorch?
Custom Code?
- Features:**
How do I explain the predictions?
What about detecting outliers and skew?
How do I wire up custom pre and post processing?

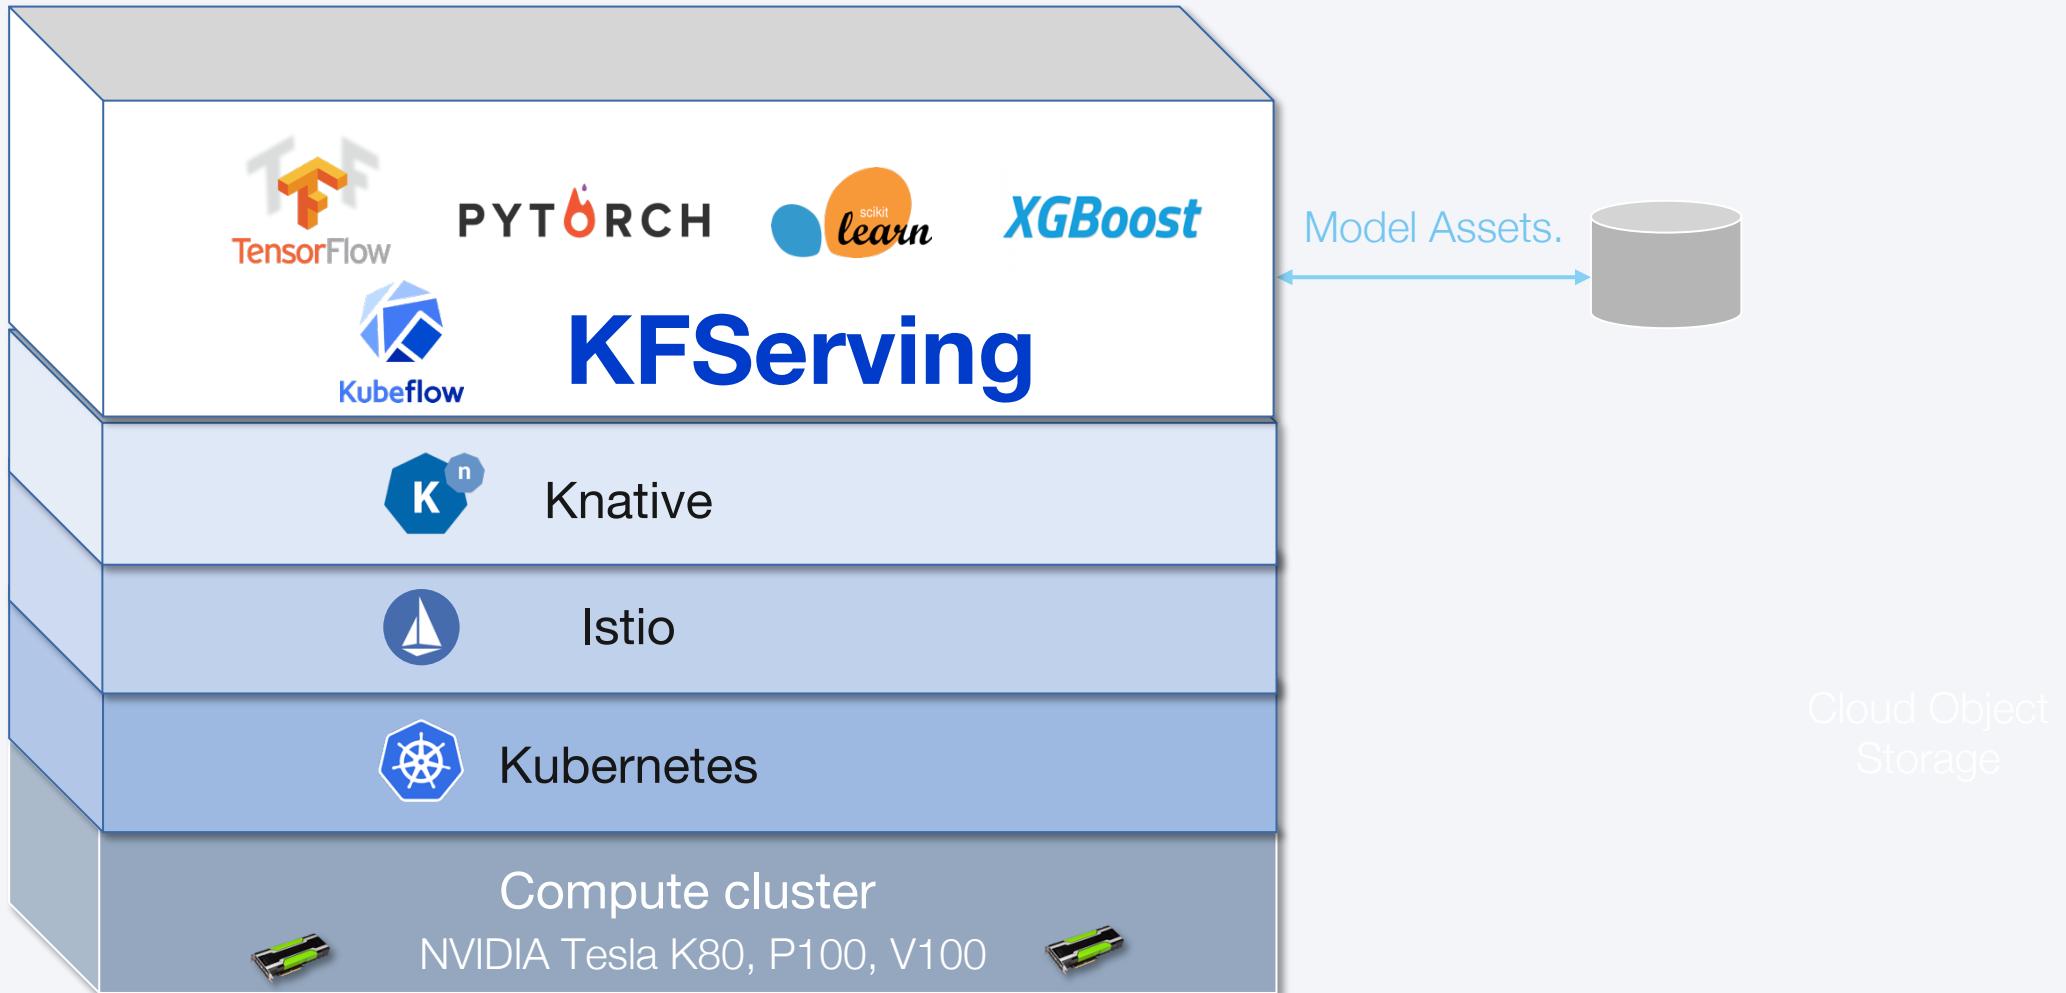
Enter: KFServing (Kubeflow Serving)



KFServing

KFServing: Model Serving and Management

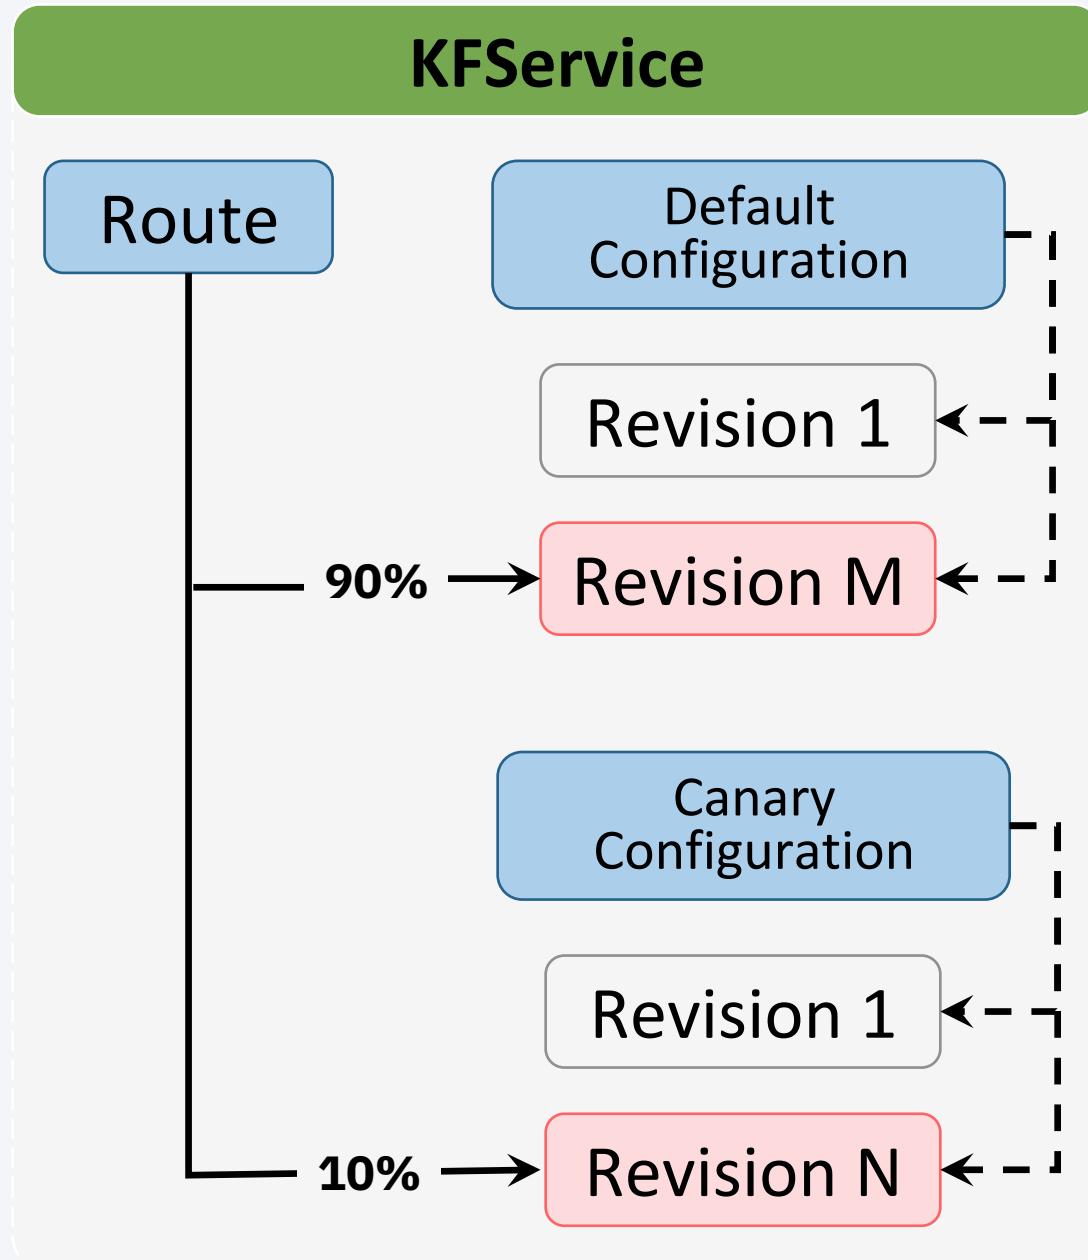
Bringing the power of Knative and Istio for serverless Model deployments



KFServing: Default and Canary Configurations

Manages the hosting aspects of your models

- **KFServe** - manages the lifecycle of models
- **Configuration** - manages history of model deployments. Two configurations for default and canary.
- **Revision** - A snapshot of your model version
 - Config and image
- **Route** - Endpoint and network traffic management



KFServing

Now (0.1.0) Support

Model Source

- s3
- gs
- pvc
- file in image

Model Framework

- tensorflow
- tensorRT
- Pytorch
- Sklearn
- XGBoost
- custom

KFServing Roadmap

Multi-Model Serving.

- Multiple KFServices share resources in the backend.
- GPU Sharing.

Flexible Inference Graphs [MLGraph CRD](#).

- Model Experimentation.
- Ensembling.
- Multi Arm Bandit.

Trusted Model Serving: Bias, Skew, and Outlier Detection.

- Online support in graph.
- Offline support with Payload Logging.

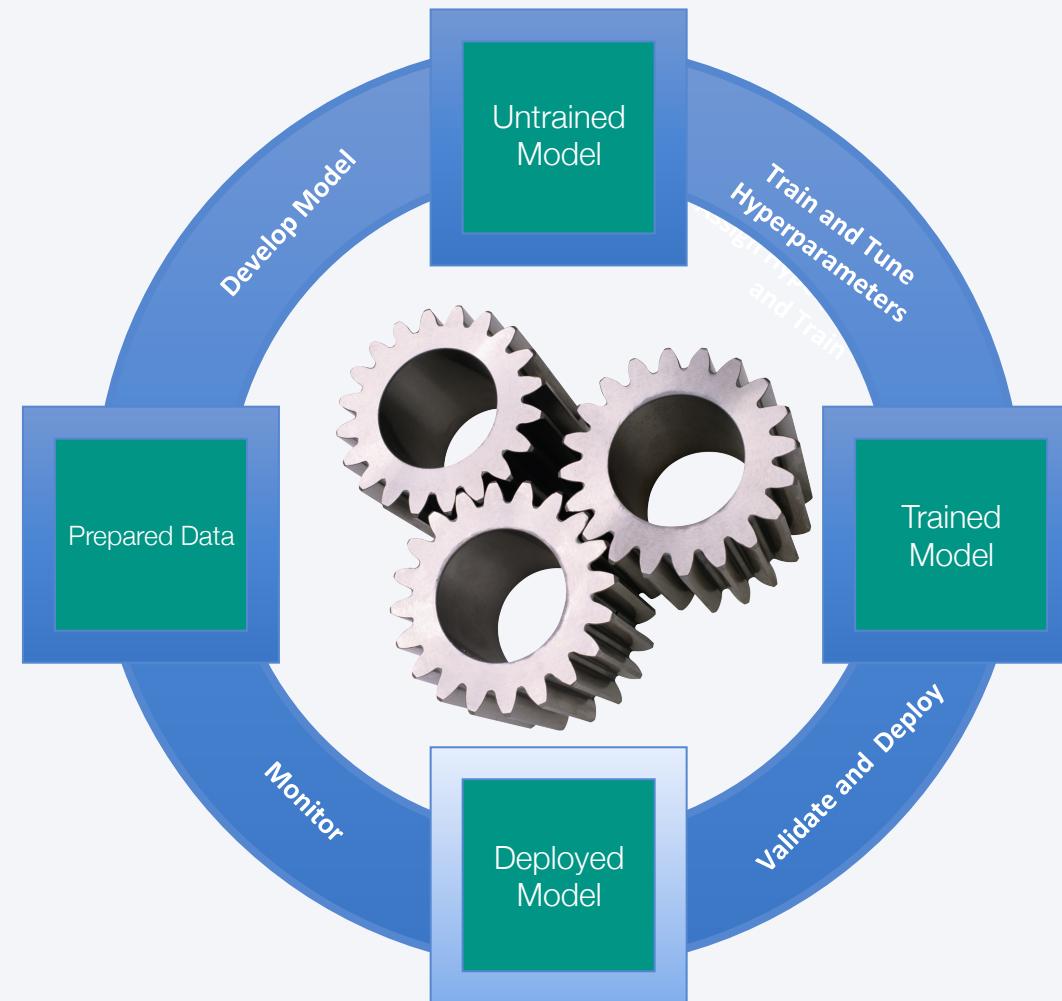
Meta-Protocol Definition.

- Unify disparate protocols across frameworks.

Adaptive batching support

- Queue and batch requests to increase throughput.

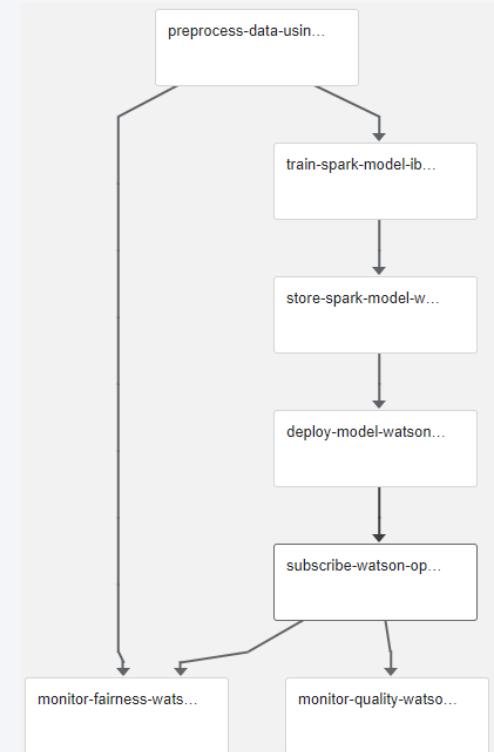
So Model is Trained, Tested and Deployed. How can we orchestrate his lifecycle together? We need AI Pipelines



AI Pipelines

Pipelines

- Covers the whole lifecycle from data and source code to the application of models
 - end-to-end experiments – effective, efficient and reliable
 - lifecycle management – data engineering, train, validate, store, deploy, serve, inference, monitor, log, ...
 - share and reuse
 - Usually represented as DAGs (Directed Acyclic Graph)
 - Blocks or steps have input and output dependencies
 - Triggers and eventing
 - Stores artifacts and metadata throughout pipeline
- ⇒ help solving “the last mile” problem**



Infrastructure to support AI Pipelines

Functions

- compose, create, orchestrate, schedule, eventing, ...
 - API and SDK
- deploy, manage, ...
 - Kubernetes or other base infrastructure
 - portable, scalable, hybrid
- monitor, logging, visualization, ...
 - UI, experiment tracking, metrics and other artifacts logging

Infrastructure to support AI Pipelines

Desired

- Easy deployment
 - One click deploy
 - One click to install components / operators (for example, on OpenShift)
 - Cloud agnostic (any Kubernetes clusters)
- Easy maintenance
 - Not to add the complexity of workflow
 - CI/CD pipelines
- Data and model versioning (b/c consistent reproducibility)
 - metadata, storage (S3, github repo, ...)
 - portable model packaging for easy serving



Infrastructure to support AI Pipelines

- More important to run on container platform
 - Container – packaging the execution environment (libraries, binaries, config) with the application abstracted in a pre-built image (eg. Docker)
 - ✓ hybrid target environments: personal laptop, private or public cloud
 - ✓ run on hosts, VMs, Kubernetes alike clusters
 - ✓ self-contained/isolation
 - ✓ lightweight than VM
 - With Kubernetes
 - ✓ container orchestration
 - ✓ scalable
 - ✓ portable
 - ✓ service health monitoring

Open source projects for Pipelines

Argo Workflows



- <https://argoproj.github.io/argo>, <https://github.com/argoproj/argo>
- container-native workflow engine for Kubernetes, each step is a container
- workflow as CRD (Custom Resource Definition, .yaml)
- can run on any Kubernetes clusters
- orchestrate highly parallel tasks
- pipelines represented as DAGs
 - compose a workflow spec in yaml, run with Argo CLI or kubectl
 - input/output (artifacts/parameters), loops, conditionals, exit handler, daemon/sidecar container, and more



```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: dag-diamond-
spec:
  entrypoint: diamond
  templates:
  - name: diamond
    dag:
      tasks:
      - name: A
        template: echo
        arguments:
          parameters: [{name: message, value: A}]
      - name: B
        dependencies: [A]
        template: echo
        arguments:
          parameters: [{name: message, value: B}]
      - name: C
        dependencies: [A]
        template: echo
        arguments:
          parameters: [{name: message, value: C}]
      - name: D
        dependencies: [B, C]
        template: echo
        arguments:
          parameters: [{name: message, value: D}]
    - name: echo
      inputs:
        parameters:
        - name: message
      container:
        image: alpine:3.7
```

Open source projects for Pipelines

Airflow

- <https://airflow.apache.org/>, <https://github.com/apache/airflow>
- A platform to programmatically author, schedule and monitor workflows
 - DAGs, Operators, Tasks, TaskInstances, ...
- Rich command line utilities and user interface
- Pipeline definitions written in Python
- KubernetesOperator, KubernetesPodOperator – each TaskInstance runs in a new pod
- (possibly used by mlflow)



```
from airflow import DAG
from airflow.operators.bash_operator import BashOperator
from datetime import datetime, timedelta

default_args = {
    'owner': 'airflow',
    'depends_on_past': False,
    'start_date': datetime(2015, 6, 1),
    'email': ['airflow@example.com'],
    'email_on_failure': False,
    'email_on_retry': False,
    'retries': 1,
    'retry_delay': timedelta(minutes=5),
    # 'queue': 'bash_queue',
    # 'pool': 'backfill',
    # 'priority_weight': 10,
    # 'end_date': datetime(2016, 1, 1),
}

dag = DAG('tutorial', default_args=default_args, schedule_interval=timedelta(days=1))

# t1, t2 and t3 are examples of tasks created by instantiating operators
t1 = BashOperator(
    task_id='print_date',
    bash_command='date',
    dag=dag)

t2 = BashOperator(
    task_id='sleep',
    bash_command='sleep 5',
    retries=3,
    dag=dag)

templated_command = """
    {% for i in range(5) %}
        echo "{{ ds }}"
        echo "{{ macros.ds_add(ds, 7) }}"
        echo "{{ params.my_param }}"
    {% endfor %}
"""

t3 = BashOperator(
    task_id='templated',
    bash_command=templated_command,
    params={'my_param': 'Parameter I passed in'},
    dag=dag)

t2.set_upstream(t1)
t3.set_upstream(t1)
```

Open source projects for Pipelines

Data Pipelines – Pachyderm



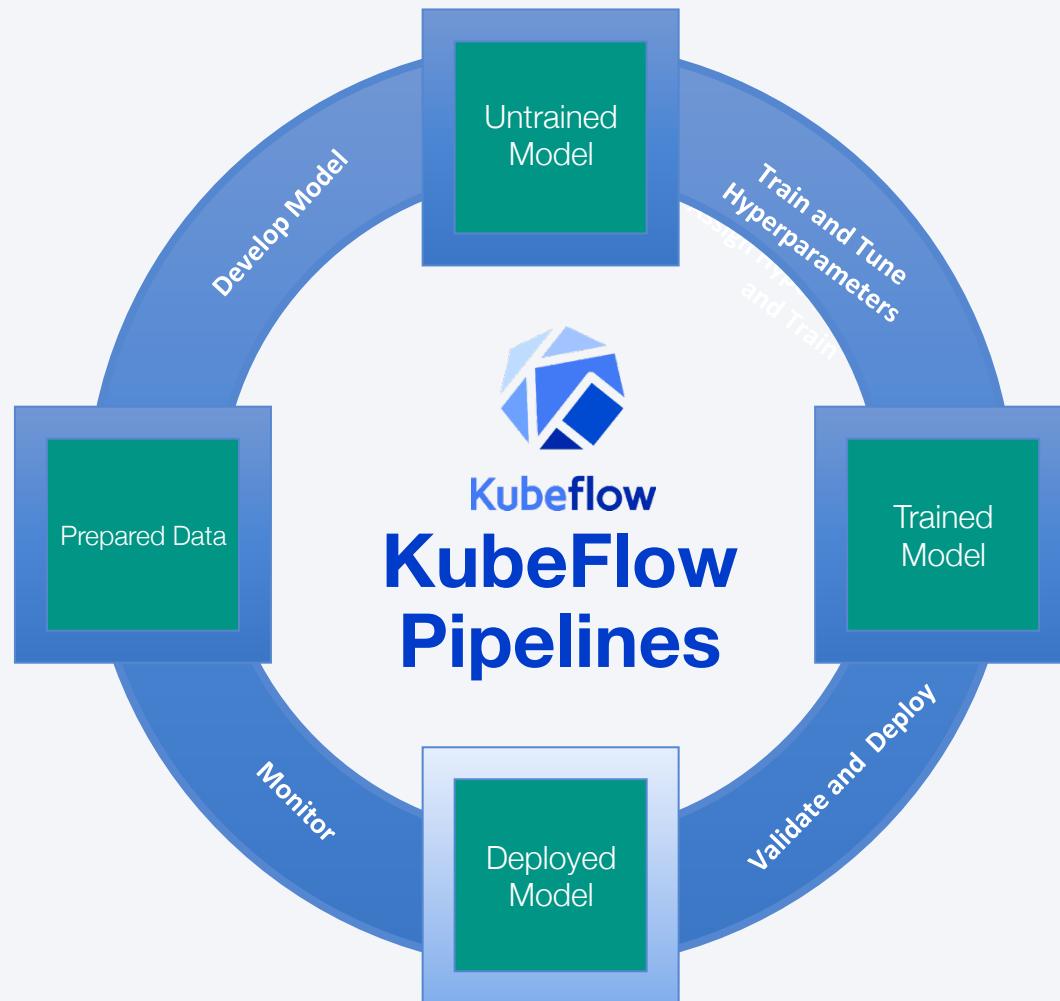
- version control for data and other artifacts
- run on container platform
- data provenance
 - track any result all the way back to its raw input data, including all analysis, code, and intermediate results
- parallelism (datum ~ worker)
- create-repo, put-file, create-pipeline
- add more data → trigger new pipelines job
- CLI, python-pachyderm client, ...
- pachyderm + kubeflow pipelines ?



```
# edges.json
{
  "pipeline": {
    "name": "edges"
  },
  "transform": {
    "cmd": [ "python3", "/edges.py" ],
    "image": "pachyderm/opencv"
  },
  "input": [
    "pfs": {
      "repo": "images",
      "glob": "*"
    }
  ]
}

# montage.json
{
  "pipeline": {
    "name": "montage"
  },
  "input": [
    {
      "cross": [
        "pfs": {
          "glob": "/",
          "repo": "images"
        }
      ],
      "pfs": {
        "glob": "/",
        "repo": "edges"
      }
    }
  ],
  "transform": {
    "cmd": [ "sh" ],
    "image": "v4tech/imagemagick",
    "stdin": [
      "montage -shadow -background SkyBlue -geometry 300x300+2+2 $(find /pfs -type f | sort) /pfs/out/montage.png"
    ]
  }
}
```

Last but not the least: Kubeflow Pipelines

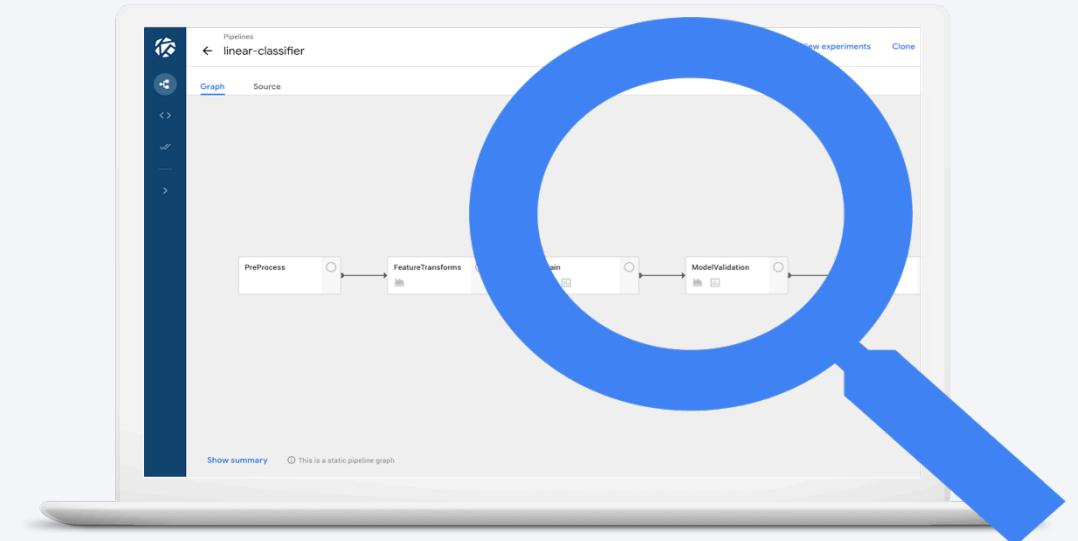


Kubeflow Pipelines

- Released to Kubeflow in Nov 2018, integrated into KF deployment CLI and 1-click-deploy-app
- Aimed to bring
 - Orchestration for complex ML workflows
 - Reproducible and reliable experimentation
 - Bridging experimentation and operationalization
 - Composition and reusable ML components and pipelines

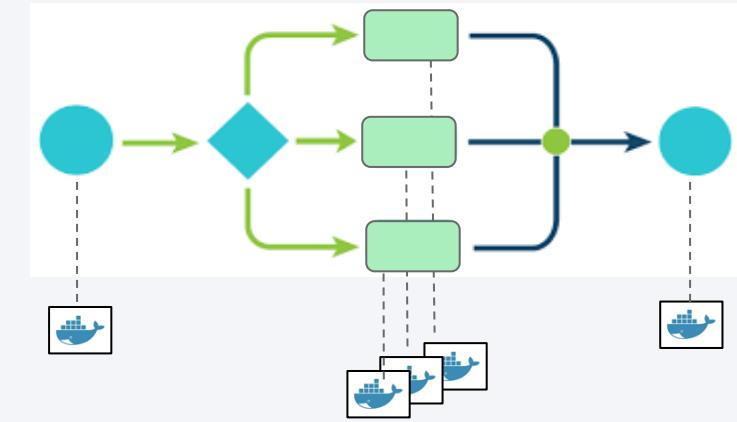


Kubeflow Pipelines



What constitutes a Kubeflow ML Pipeline

- Containerized implementations of ML Tasks
 - Pre-built components: Just provide params or code snippets (e.g. training code)
 - Create your own components from code or libraries
 - Use any runtime, framework, data types
 - Attach k8s objects - volumes, secrets
- Specification of the sequence of steps
 - Specified via Python DSL
 - Inferred from data dependencies on input/output
- Input Parameters
 - A “Run” = Pipeline invoked w/ specific parameters
 - Can be cloned with different parameters
- Schedules
 - Invoke a single run or create a recurring scheduled pipeline



Experiments > My experiment
← My first run

Graph Run output Config

Artifacts Input/Output Logs

Input parameters

url1 gs://ml-pipeline-playground/shakespeare1.txt

Output parameters

download1-downloaded

With which he yoketh your rebellious necks Razeth your cities and subverts your towns And in a moment makes them desolate

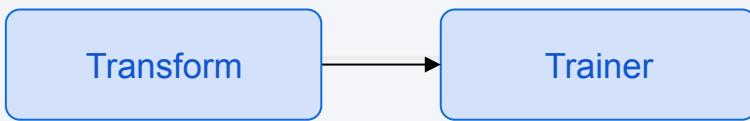
Pipeline name	Description	Uploaded on
[Sample] Basic - Condition	A pipeline shows how to use dsl.Condition. For source code, refer to https://github.com/kubeflow/pipelines/tree/master/samples/by-example/basic/dsl.Condition.go	02/01/2019, 11:24:37
[Sample] Basic - Exit Handler	A pipeline that downloads a message and print it out. Exit Handler will run at the end. For source code, refer to https://github.com/kubeflow/pipelines/tree/master/samples/by-example/basic/dsl.ExitHandler.go	02/01/2019, 11:24:36
[Sample] Basic - Immediate...	A pipeline with parameter values hard coded. For source code, refer to https://github.com/kubeflow/pipelines/tree/master/samples/by-example/basic/dsl.Immediate.go	02/01/2019, 11:24:34
[Sample] Basic - Parallel Join	A pipeline that downloads two messages in parallel and print the concatenated result. For source code, refer to https://github.com/kubeflow/pipelines/tree/master/samples/by-example/basic/dsl.ParallelJoin.go	02/01/2019, 11:24:33
[Sample] Basic - Sequential	A pipeline with two sequential steps. For source code, refer to https://github.com/kubeflow/pipelines/tree/master/samples/by-example/basic/dsl.Sequential.go	02/01/2019, 11:24:32
[Sample] ML - TFX - Taxi Tip...	Example pipeline that does classification with model analysis based on a public tax cab BI...	02/01/2019, 11:24:30
[Sample] ML - XGBoost - Trai...	A trainer that does end-to-end distributed training for XGBoost models. For source code, re...	02/01/2019, 11:24:29

+ Upload pipeline Refresh Delete

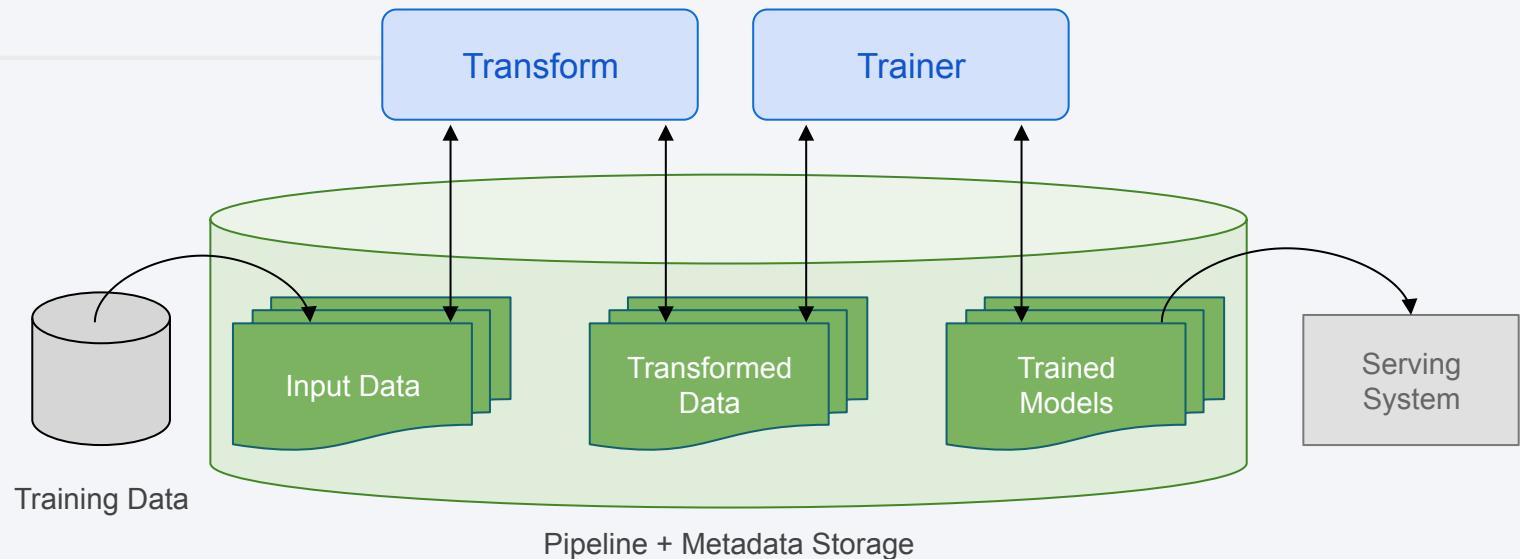
Rows per page: 10 < >

Metadata and Pipeline Execution

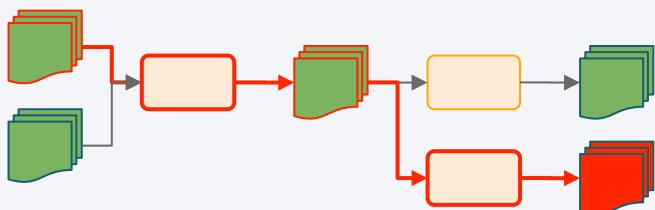
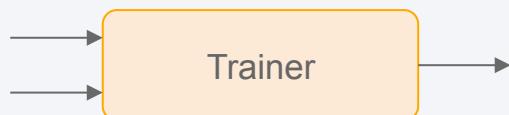
Task-Aware Pipelines



Task- and Data-Aware Pipelines



What's in the Metadata Store?



Type definitions of Artifacts and their Properties
Models, Data, Evaluation Metrics, Visualizations
(core types come out of the box, custom types can be defined)

Execution Records (Runs)

Runtime configuration params, Inputs + Outputs references, versioning of binaries/images

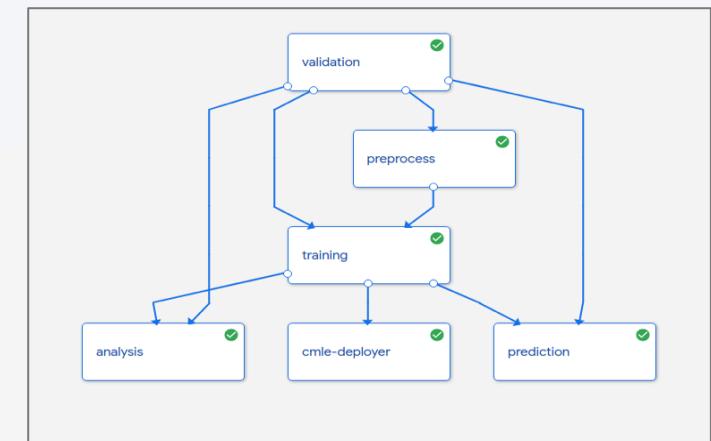
Lineage Tracking Across All Executions

To recurse back to all inputs of a specific artifact.
Find out how a given dataset was used? What was used to create a given model?

Define Pipeline with Python SDK

```
@dsl.pipeline(name='Taxi Cab Classification Pipeline Example')
def taxi_cab_classification(
    output_dir,
    project,
    Train_data      = 'gs://bucket/train.csv',
    Evaluation_data = 'gs://bucket/eval.csv',
    Target          = 'tips',
    Learning_rate   = 0.1, hidden_layer_size = '100,50', steps=3000):

    tfdv           = TfdvOp(train_data, evaluation_data, project, output_dir)
    preprocess     = PreprocessOp(train_data, evaluation_data, tfdv.output["schema"], project, output_dir)
    training       = DnnTrainerOp(preprocess.output, tfdv.schema, learning_rate, hidden_layer_size, steps,
                                target, output_dir)
    tfma           = TfmaOp(training.output, evaluation_data, tfdv.schema, project, output_dir)
    deploy         = TfServingDeployerOp(training.output)
```



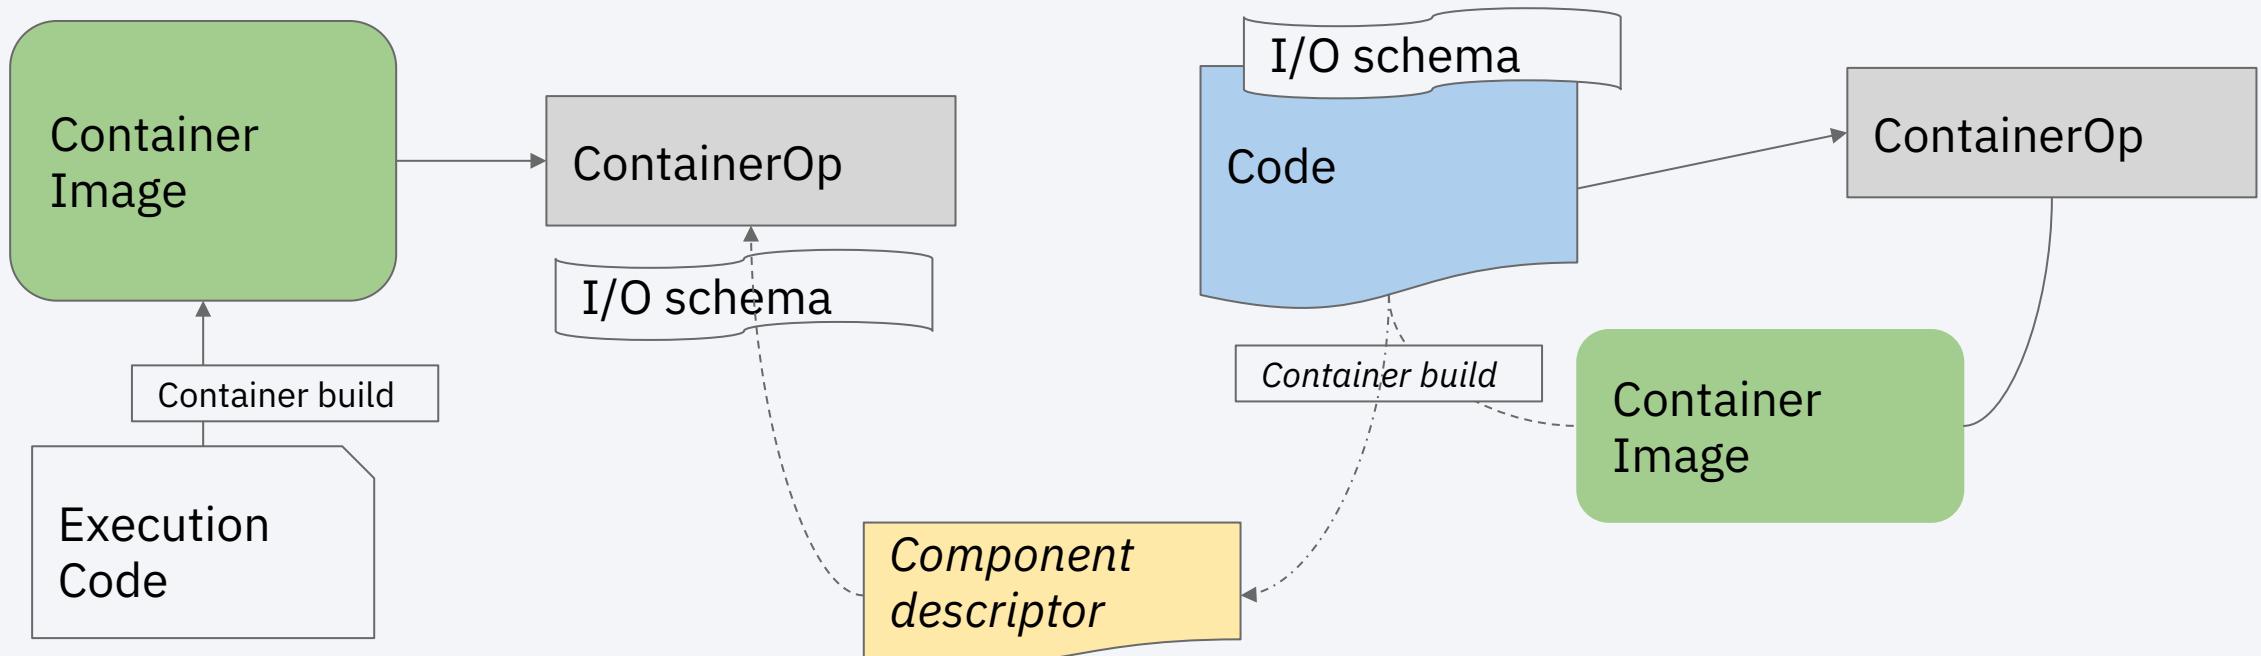
Compile and Submit Pipeline Run

```
dsl.compile(taxi_cab_classification, 'tfx.tar.gz')
run = client.run_pipeline(
    'tfx_run', 'tfx.tar.gz', params={'output': 'gs://dpa22', 'project': 'my-project-33'})
```

Creating your own components

Ways to build reusable components for pipelines

- Create a container with your code and write either a ContainerOp() or shareable component descriptor
- Turn your python code into a component directly in the notebook (with or without building a container)
 - These components can be exported into a shareable format



Simply Create Component from a Python function..

```
@dsl.python_component(  
    name='cmle_deployer',  
    description='deploys a model to GCP CMLE',  
    base_image='python:3.5')  
def deploy_model(model_dot_version: str, model_path: str, gcp_project: str):  
    # code to deploy a model  
    ...  
    ...
```

Locally test it and convert to a component with a container image

```
MyDeployerOp = compiler.build_python_component(  
    component_func=deploy_model,  
    target_image=gcr.io/myproject33/mydeployer-image:1.0')
```

Use it in a Pipeline

```
@dsl.pipeline(  
    name='Taxi Cab Classification Pipeline Example',  
)  
def taxi_cab_classification():  
  
    ...  
    deploy = TfServingDeployerOp(training.output) ⇒ deploy = MyDeployerOp(name, training.output, project)
```

Create Component from a Python function without building container

```
In [ ]: import kfp.components as comp
```

Simple function that just add two numbers:

```
In [ ]: #Define a Python function
def add(a: float, b: float) -> float:
    '''Calculates sum of two arguments'''
    return a + b
```

Convert the function to a pipeline operation

```
In [ ]: add_op = comp.func_to_container_op(add)
```

Can export it to a file with component definition from here

Use it in a Pipeline

```
In [ ]: import kfp.dsl as dsl
@dsl.pipeline(
    name='Calculation pipeline',
    description='A toy pipeline that performs arithmetic calculations.'
)
def calc_pipeline(
    a='a',
    b='7',
    c='17',
):
    #Passing pipeline parameter and a constant value as operation arguments
    add_task = add_op(a, 4) #Returns a dsl.ContainerOp class instance.

    #Passing a task output reference as operation arguments
    #For an operation with a single return value, the output reference can be accessed using `task.output` or `task.outputs['output_name']` syntax
    divmod_task = divmod_op(add_task.output, b, '/')

    #For an operation with a multiple return values, the output references can be accessed using `task.outputs['output_name']` syntax
    result_task = add_op(divmod_task.outputs['quotient'], c)
```

Pipelines

Pipelines for your machine learning workloads.

[VIEW ALL PIPELINES](#)

[UPLOAD A PIPELINE](#)

Train with FfDL and canary testing with Knative

This pipeline trains 2 models, does fairness and robustness checking, then does canary test with Knative for deployment.

[Open Source](#)

Train and deploy with FfDL and Seldon (OSS)

A simple IBM OSS pipeline demonstrates how to train a model using Fabric for Deep Learning and then deploy it with Seldon.

[Open Source](#)

Train and deploy with Watson Machine Learning

A sample pipeline for training, storing and deploying a Tensorflow model with MNIST handwriting recognition on IBM Watson Machine Learning service.

[IBM Watson](#)

Model analysis and operations using AI OpenScale

IBM AI OpenScale to monitor and analyze ML models.

[Open Source](#)

End-to-end pipeline with fairness and robustness test.

This AI pipeline uses FfDL to train a model, does fairness checking and robustness checking. For successful checking results, it does the model deployment with A/B testing.

[IBM Watson](#)

Kubernetes model deployment

Deploy a container to a Kubernetes platform with its own service.

[IBM Watson](#)

Knative model deployment

Deploy a container to a Knative platform with traffic splitting.

[Open Source](#)

[Sample] Basic - Condition

A pipeline shows how to use dsl.Condition. For source code, refer to <https://github.com/kubeflow/pipelines/blob/master/samples/basic/condition.py>.

[IBM Watson](#)

[Sample] Basic - Parallel Join

A pipeline that downloads two messages in parallel and print the concatenated result. For source code, refer to https://github.com/kubeflow/pipelines/blob/master/samples/basic/parallel_join.py.

[IBM Watson](#)

[Sample] ML - XGBoost - Training with Confusion Matrix

A trainer that does end-to-end distributed training for XGBoost models. For source code, refer to <https://github.com/kubeflow/pipelines/tree/master/samples/xgboost-spark>.

[Open Source](#)

Pipelines

Train a model, do fairness and robustness checks, and deploy it with A/B testing.

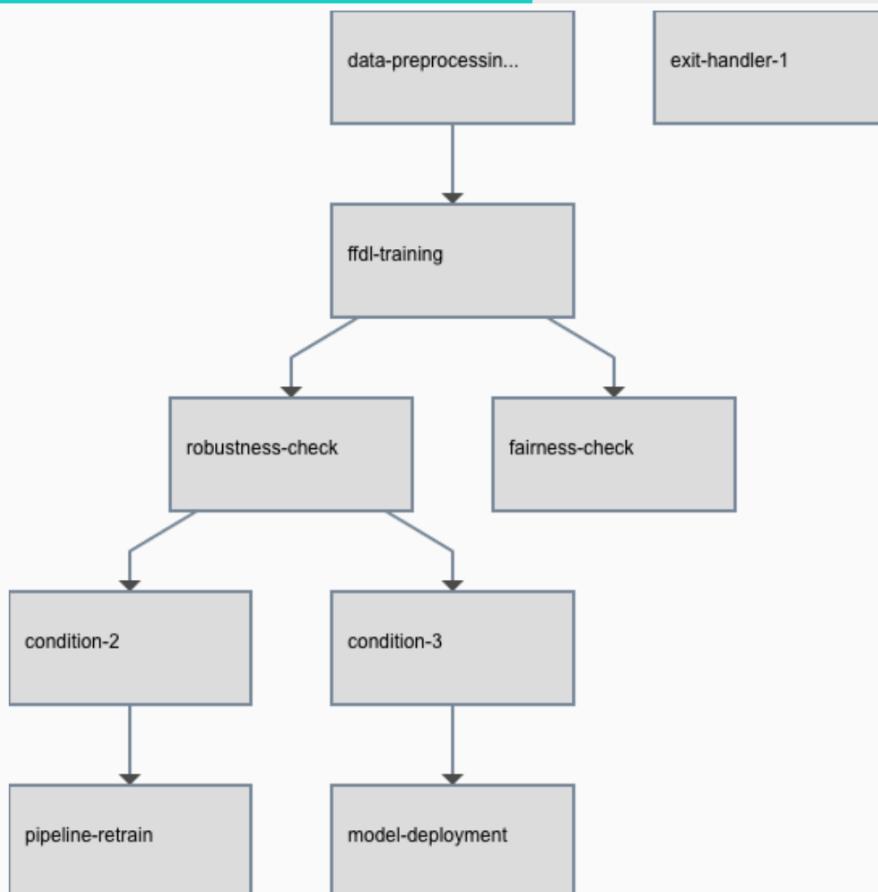
← PIPELINES

GET CODE

GRAPH

DETAILS

PIPELINE YAML



```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: ffdl-pipeline-
spec:
  arguments:
    parameters:
      - name: model-def-file-path
        value: gender_classification_training.zip
      - name: manifest-file-path
        value: manifest_training.yml
      - name: preprocessing-def-file-path
        value: preprocessing.zip
      - name: preprocessing-manifest-file-path
        value: preprocessing_manifest.yml
      - name: fgsm-attack-epsilon
        value: '0.2'
      - name: retrain-pipe-exp-id
        value: 5adacf87-48aa-4c6e-9aad-9a9cf933b9a6
      - name: retrain-pipeline-id
        value: f246bd8a-832e-4560-a8db-3cbee851c3fd
      - name: pipeline-url
        value: http://169.48.196.99/pipeline
      - name: primary-model-revision
        value: model-serving-00001
      - name: new-model-traffic-percentage
        value: '5'
      - name: model-class-file
        value: gender_classification_training.py
      - name: model-class-name
        value: ThreeLayerCNN
```

Components

Components that can be used to build your pipelines.

[VIEW ALL COMPONENTS](#)[UPLOAD A COMPONENT](#)

Deploy Model - Watson Machine Learning

Deploy stored model on Watson Machine Learning as a web service.

[IBM Watson Machine Learning](#)

Subscribe - Watson OpenScale

Binding deployed models and subscribe them to Watson OpenScale service.

[IBM Watson OpenScale](#)

Store model - Watson Machine Learning

Store and persistent trained model on Watson Machine Learning.

[IBM Watson Machine Learning](#)

Knative model deploy

Deploy AI models using Knative serving.

[OpenSource](#)

Fabric for Deep Learning - Train Model

Train Machine Learning and Deep Learning Models remotely using Fabric for Deep Learning

[OpenSource](#)

Model Robustness Check

Perform a robustness check using fast gradient attack with ART to make sure the model is robust against simple gradient changes.

[OpenSource](#)

Kubernetes model deploy

Deploy AI models using Kubernetes deployment.

[OpenSource](#)

Train Model - Watson Machine Learning

Train Machine Learning and Deep Learning Models in the Cloud using Watson Machine Learning

[IBM Watson Machine Learning](#)

Seldon Core - Serve PyTorch Model

Serve PyTorch Models remotely as web service using Seldon Core

[OpenSource](#)

Model Fairness Check

Perform a fairness check on a certain attribute using AIF360 to make sure the model is fair and ethical

[OpenSource](#)

Models

Machine learning models that can be used in your pipelines.

[VIEW ALL MODELS](#)[UPLOAD A MODEL](#)

MAX Audio Classifier

IBM Model Asset eXchange(MAX) model that identifies sounds in short audio clips

[Audio Classification](#)

MAX Named Entity Tagger

IBM Model Asset eXchange(MAX) model that locates and tags named entities in text

[Natural Language Processing](#)

MAX Breast Cancer Mitosis Detector

IBM Model Asset eXchange(MAX) model that detects whether a mitosis exists in an image of breast cancer tumor cells

[Image Classification](#)

MAX Facial Age Estimator

IBM Model Asset eXchange(MAX) model that recognizes faces in an image and estimate the age of each face

[Facial Recognition](#)

MAX Image Caption Generator

IBM Model Asset eXchange(MAX) model that generates captions from a fixed vocabulary describing the contents of images in the COCO dataset

[Image-To-Text Translation](#)

MAX Image Segmenter

IBM Model Asset eXchange(MAX) model that identifies objects in an image, additionally assigning each pixel of the image to a particular object

[Object Detection in image](#)

MAX News Text Generator

IBM Model Asset eXchange(MAX) model that generates English-language text similar to the news articles in the One Billion Words data set

[Language Modeling](#)

MAX Object Detector

IBM Model Asset eXchange(MAX) model that localizes and identifies multiple objects in a single image

[Object detection in images](#)

MAX Sports Video Classifier

IBM Model Asset eXchange(MAX) model that classifies sporting activities in videos

[Video Classification](#)

MAX Weather Forecaster

IBM Model Asset eXchange(MAX) model that predicts hourly weather features given historical data for a specific location

[Time Series Prediction](#)

MAX Facial Age Estimator

IBM Model Asset eXchange(MAX) model that recognizes faces in an image and estimate the age of each face

← MODELS

GET CODE

DETAILS

CREATE RUN

YAML DEFINITION

SAMPLE SERVING CODE

About:

model identifier

max-facial-age-estimator

domain

Facial Recognition

description

IBM Model Asset eXchange(MAX) model that recognizes faces in an image and estimate the age of each face

Usage Info:

framework name

Keras

v2.2.2

license

Apache 2.0

website

<https://developer.ibm.com/exchanges/models/all/max-facial-age-estimator>

Serving Details:

tested platforms

kubernetes, knative

container image

codait/max-facial-age-estimator:latest

KUBERNETES

```
asai.pipeline(  
    name='Model Deployment pipeline',  
    description='A pipeline for ML/DL model deployment on Kubernetes.'  
)  
def model_pipeline(model_id='max-facial-age-estimator'):  
    """A pipeline for ML/DL model deployment."""  
  
    from kfp import dsl  
    from ai_pipeline_params import use_ai_pipeline_params  
  
    secret_name = 'e2e-creds'  
  
    model_config = dsl.ContainerOp(  
        name='model_config',  
        image='tomcli/model-config',  
        command=['python'],  
        arguments=[  
            '-u', 'model-config.py',  
            '--secret_name', secret_name,  
            '--model_id', model_id  
        ],  
        file_outputs={  
            'model_serving_image': '/tmp/model_serving_image',  
            'deployment_name': '/tmp/deployment_name'  
        }  
    )  
  
    model_deployment = dsl.ContainerOp(
```

Notebooks

Notebooks for your data science tasks.

[VIEW ALL NOTEBOOKS](#)[UPLOAD A NOTEBOOK](#)

AIF360 Gender Classification

Notebook to train AIF360 Gender Classification with reweighing.

[OpenSource](#)

ART detector model

Notebook to train ART detector model to detect possible adversarial attack.

[OpenSource](#)

ART poisoning attack

Use Notebook to leverage ART for poisoning training data and learn how to defend it.

[OpenSource](#)

ART with FfDL

Notebook to do Adversarial training using Fabric for Deep Learning.

[OpenSource](#)

Train with FfDL and canary testing with Knative.

This notebook trains a model, does fairness and robustness checking, then does canary test with Knative for deployment.

[OpenSource](#)

End-to-end with fairness and robustness test

This notebook uses FfDL to train a model, does fairness checking and robustness checking.

[OpenSource](#)

Watson OpenScale

A notebook showcases how to train custom Spark model on the Cloud and use IBM WatsonOpenScale to monitor and analyze ML model.

[Watson OpenScale](#)

Watson OpenScale walk through

End to End Watson OpenScale walk through.

[Watson OpenScale](#)

ART poisoning attack

Use Notebook to leverage ART for poisoning training data and learn how to defends it

← NOTEBOOKS

GET CODE

DETAILS

CREATE RUN

YAML DEFINITION

NOTEBOOK CODE

About:

name ART poisoning attack

description Use Notebook to leverage ART for poisoning training data and learn how to defend it

platform OpenSource

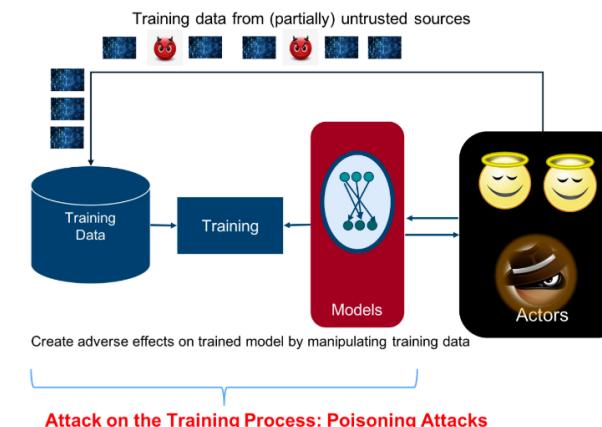
Implementation:

github https://github.com/IBM/adversarial-robustness-toolbox/blob/master/notebooks/mnist_poisoning_demo.ipynb



adversarial-robustness-toolbox / notebooks

Poisoning Attacks



Backdoor attacks:

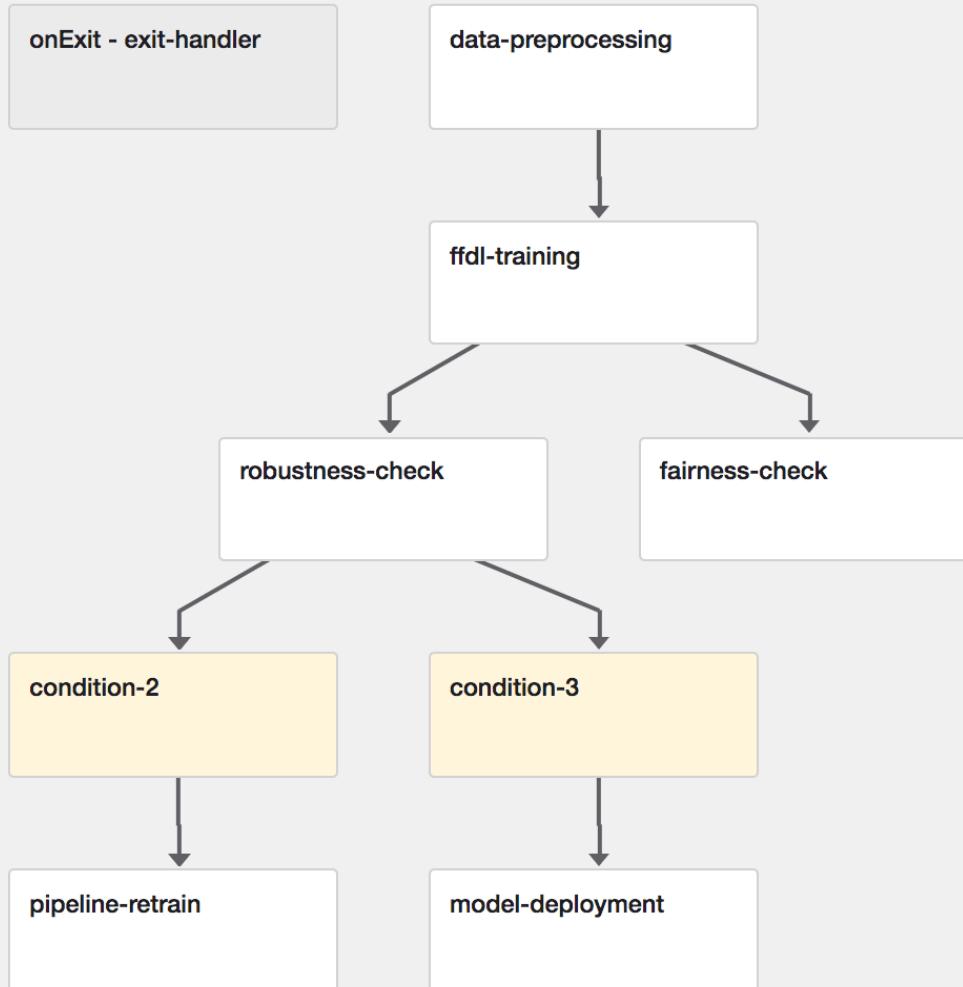
Recent work (Gu et al. 2017; Liu et al. 2017; Liu et al. 2017b) demonstrated that adversaries can generate "backdoors" into neural networks by poisoning the training data

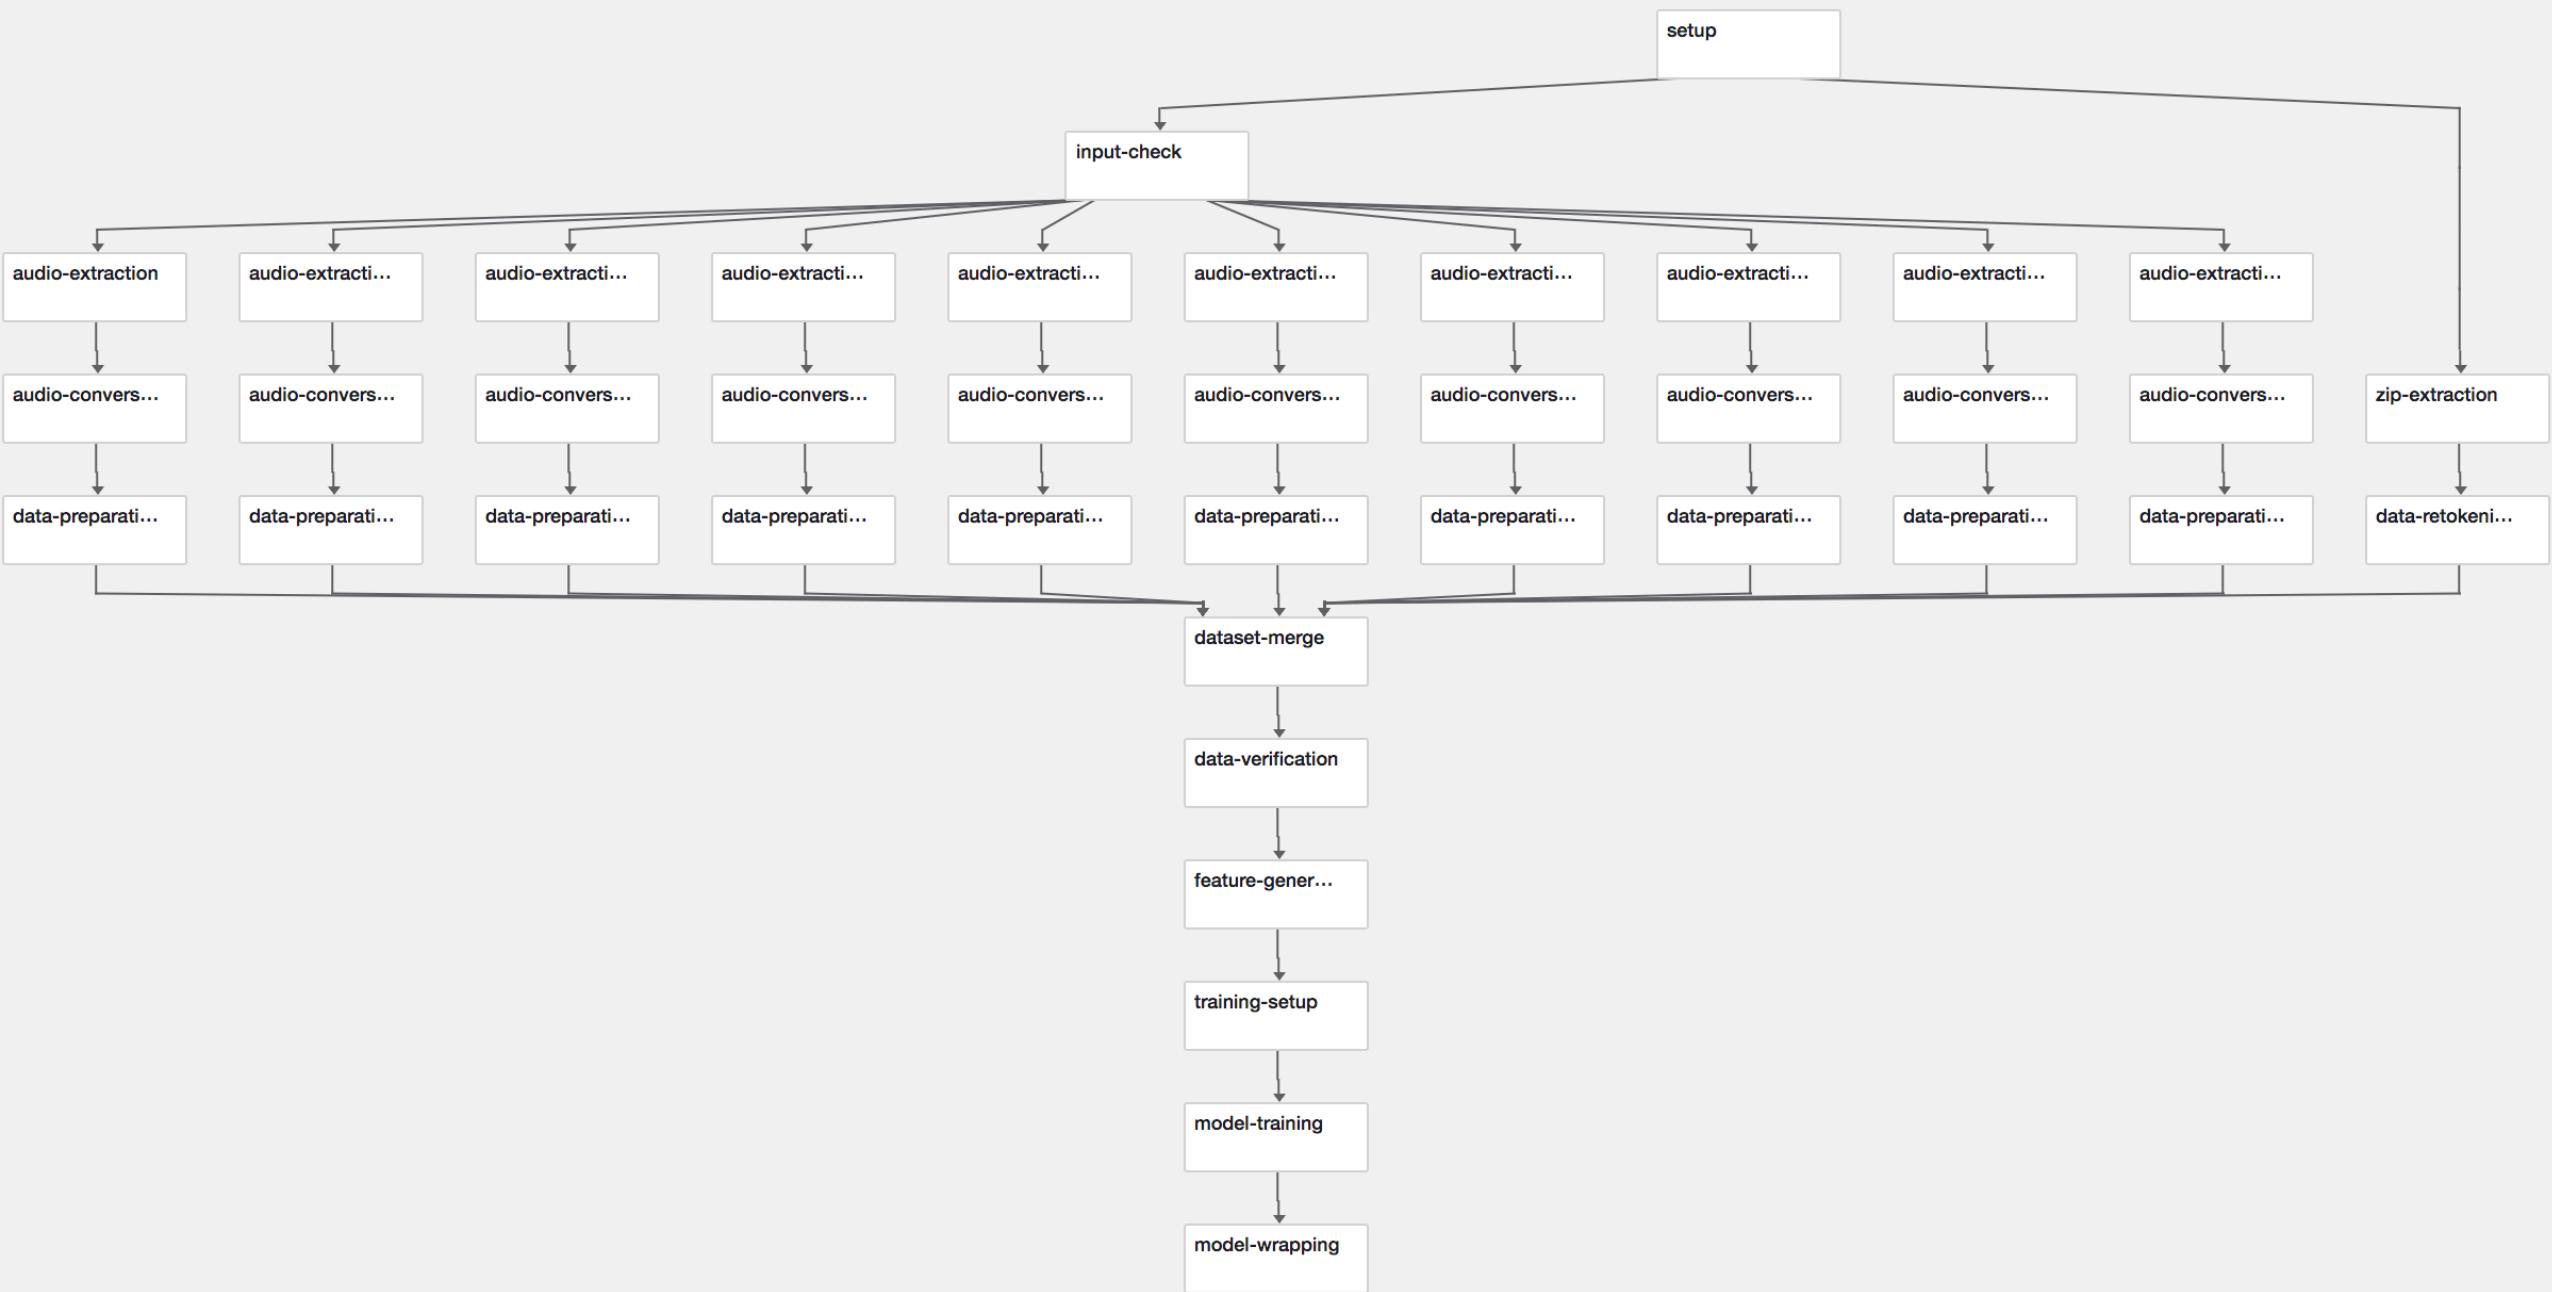
- Models with backdoors perform well on standard training and validation samples
- but behaves badly on specific attacker-chosen inputs

In [1]:

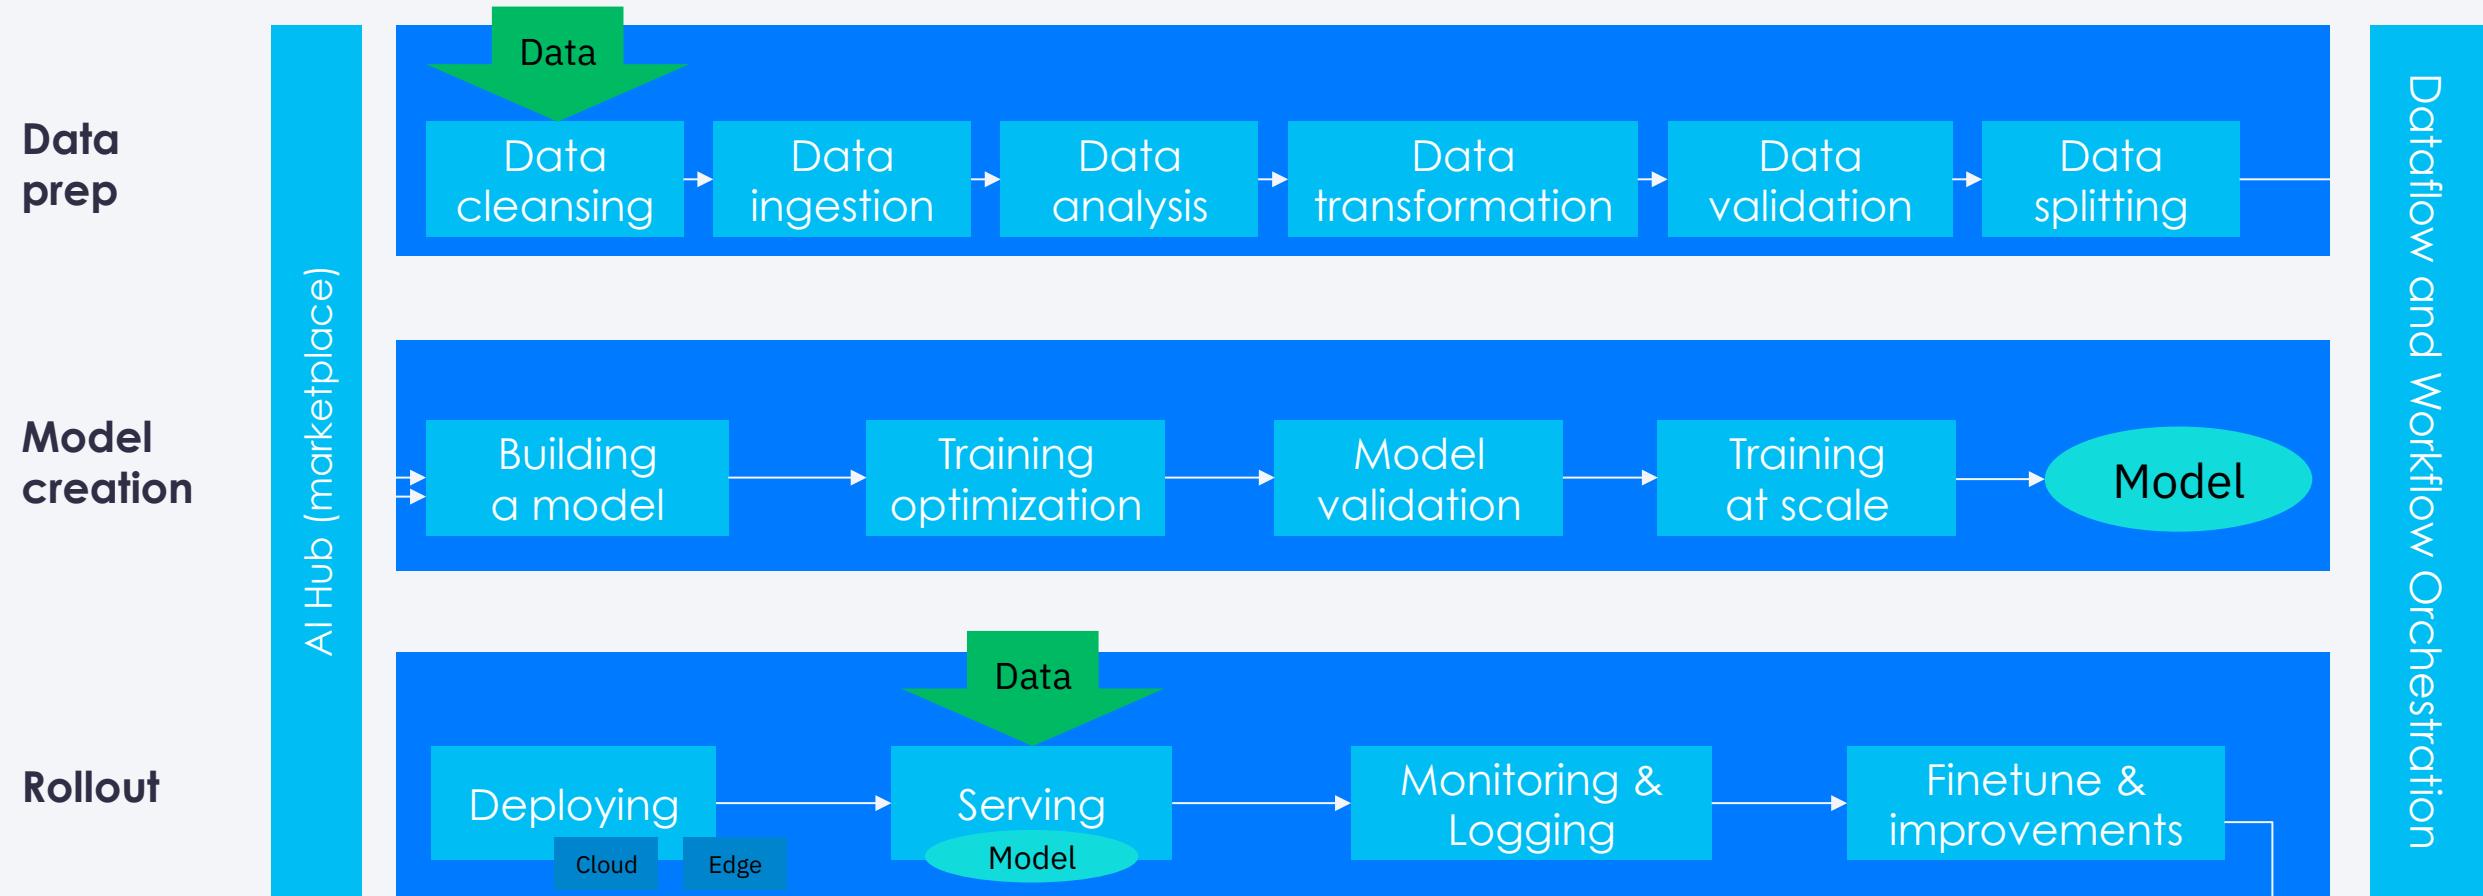
```
from __future__ import absolute_import, division, print_function, u
import os, sys
from os.path import abspath

module_path = os.path.abspath(os.path.join('..'))
```

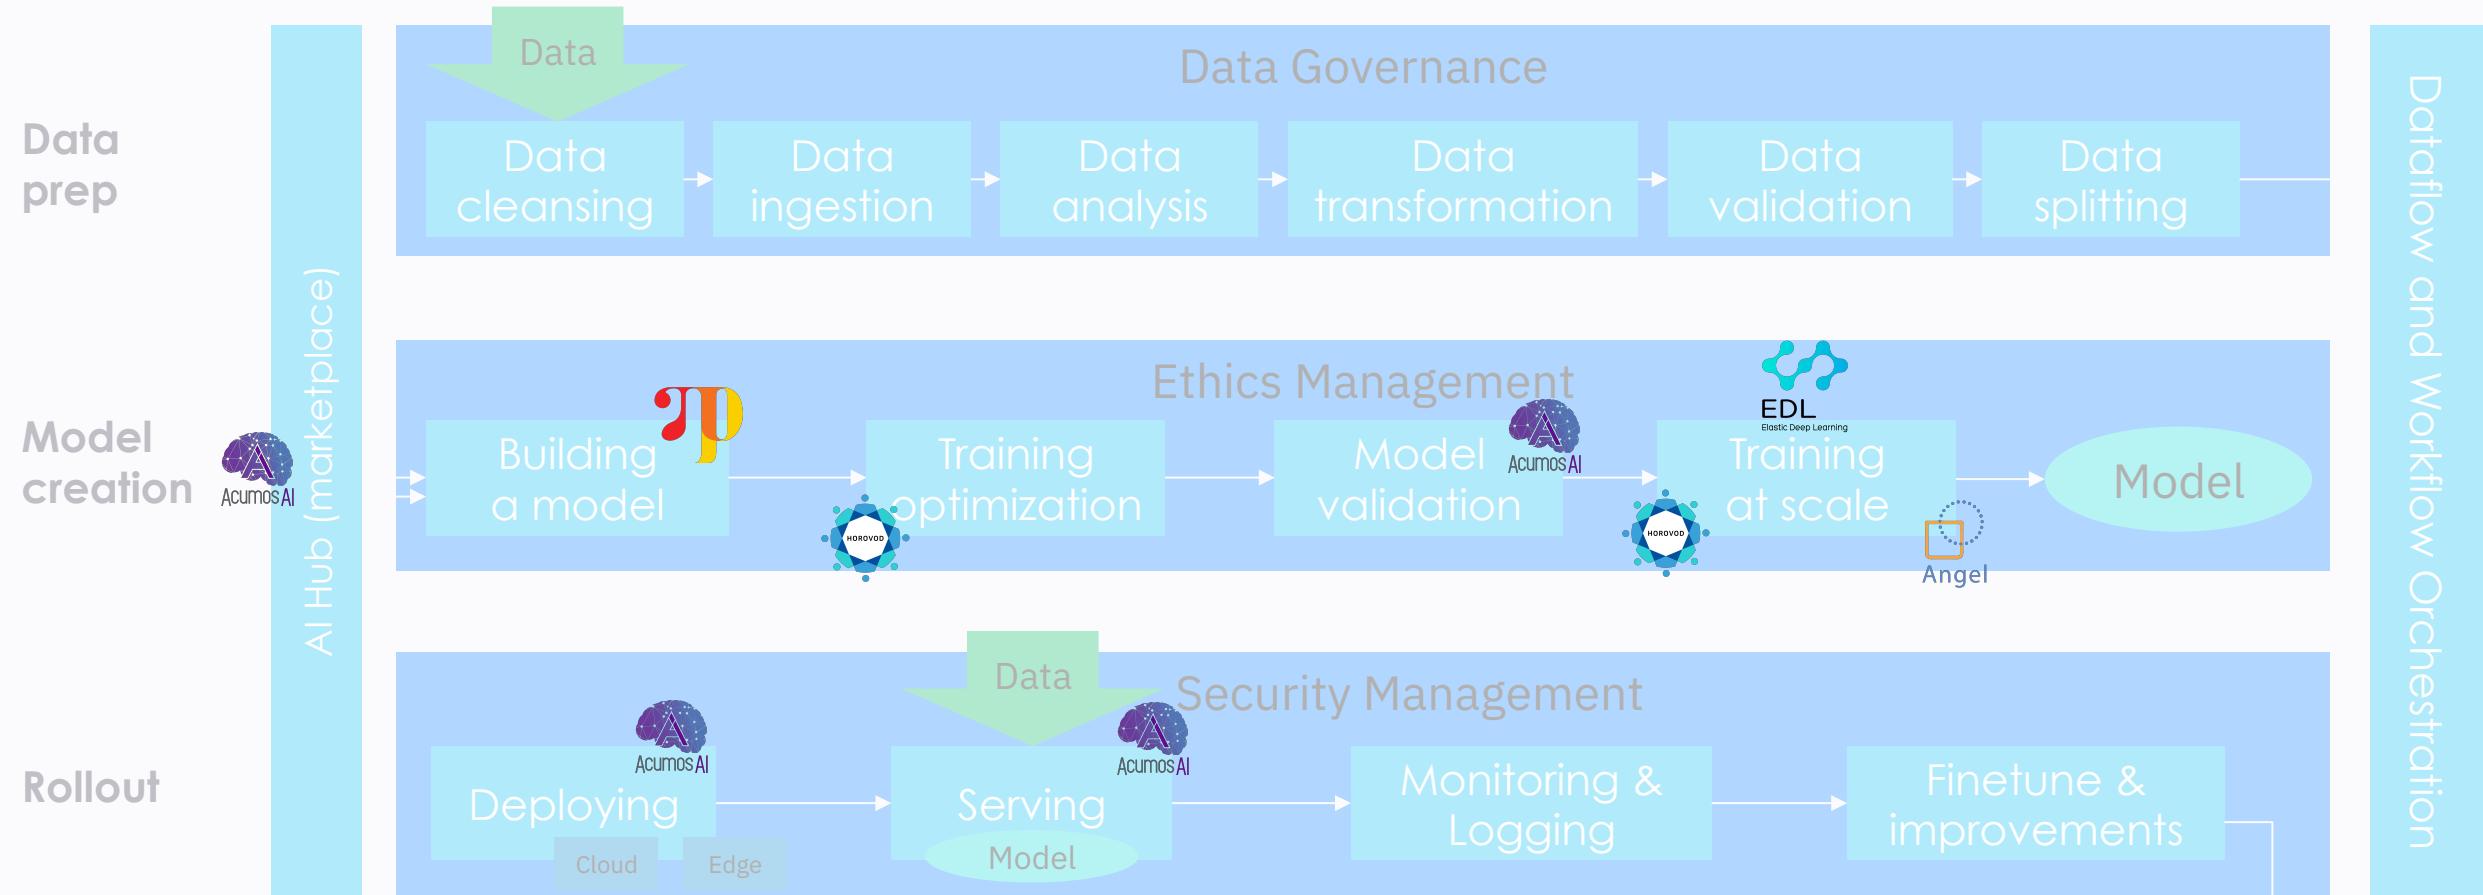
[← End-to-end pipeline with fairness and robustness test](#)[Graph](#)[Source](#)



Machine Learning Stack



LFAI Machine Learning Stack



Open Source Machine Learning Stack

